


# Chapter 8: Memory- Management Strategies

---





# Chapter 8: Memory Management Strategies

---

## OBJECTIVES

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation

## OUTLINE

- 8.1 Background
- 8.2 Swapping
- 8.3 Contiguous Memory Allocation
- 8.4 Segmentation
- 8.5 Paging





# 1. BACKGROUND

## Overview

- A **program** must be brought (from disk) into **memory** and placed within a process for it to be run.
- A program can be written in machine language, assembly language, or high-level language.
- **Main memory** and **registers** are the only storage entities that a CPU can access directly.
- The CPU fetches instructions from main memory according to the value of the **program counter**.
- Typical instruction execution cycle – fetch instruction from memory, decode the instruction, operand fetch, possible storage of result in memory.





- **Memory unit** only sees a stream of one of the following:

- **address + read** requests (e.g., load memory location 20010 into register number 8).
- **address + data** and **write** requests (e.g., store content of register 6 into memory location 1090).

Memory unit does not know how these addresses were generated.

**Register access**  
can be done in one CPU clock (or less)





## Address binding

- A **program** residing on the disk needs to be brought into **memory** in order to execute. Such a program is usually stored as a binary executable file and is kept in an **input queue**.
- In general, we do not know in advance where the program is going to reside in **memory**.

Therefore, it is convenient to assume that the first **physical address** of a program always starts at location **0000**.

- Without some hardware or software support, program must be loaded into address **0000**.

- It is impractical to have first physical address of user process to always start at location 0000.
- Most (all) computer systems provide hardware and/or software support for memory management.





- In general, addresses are represented in different ways at different stages of a program's life:

- Addresses in the source program are generally **symbolic**.
  - *i.e.*, variable "count"
- A compiler typically **binds** these symbolic addresses to **relocatable addresses**.
  - *i.e.*, "14 bytes from beginning of this module"
- Linker or loader will **bind** relocatable addresses to **absolute addresses**.
  - *i.e.*, 74014

Each **binding** maps one address space to another address space.





## Logical vs physical address space

- An address generated by CPU is commonly referred to as a **logical address** (also referred to as **virtual address**)
- Once loaded into the memory-address register of the memory is commonly referred to as a **physical address**.
- The compile-time and load-time address-binding methods generate identical **logical addresses** and **physical addresses**.

**Logical address space**  
→ the set of all **logical addresses** generated by a program.


**Physical address space**  
→ the set of all **physical addresses** corresponding to the **logical addresses**.





# Memory-Management Unit (MMU)

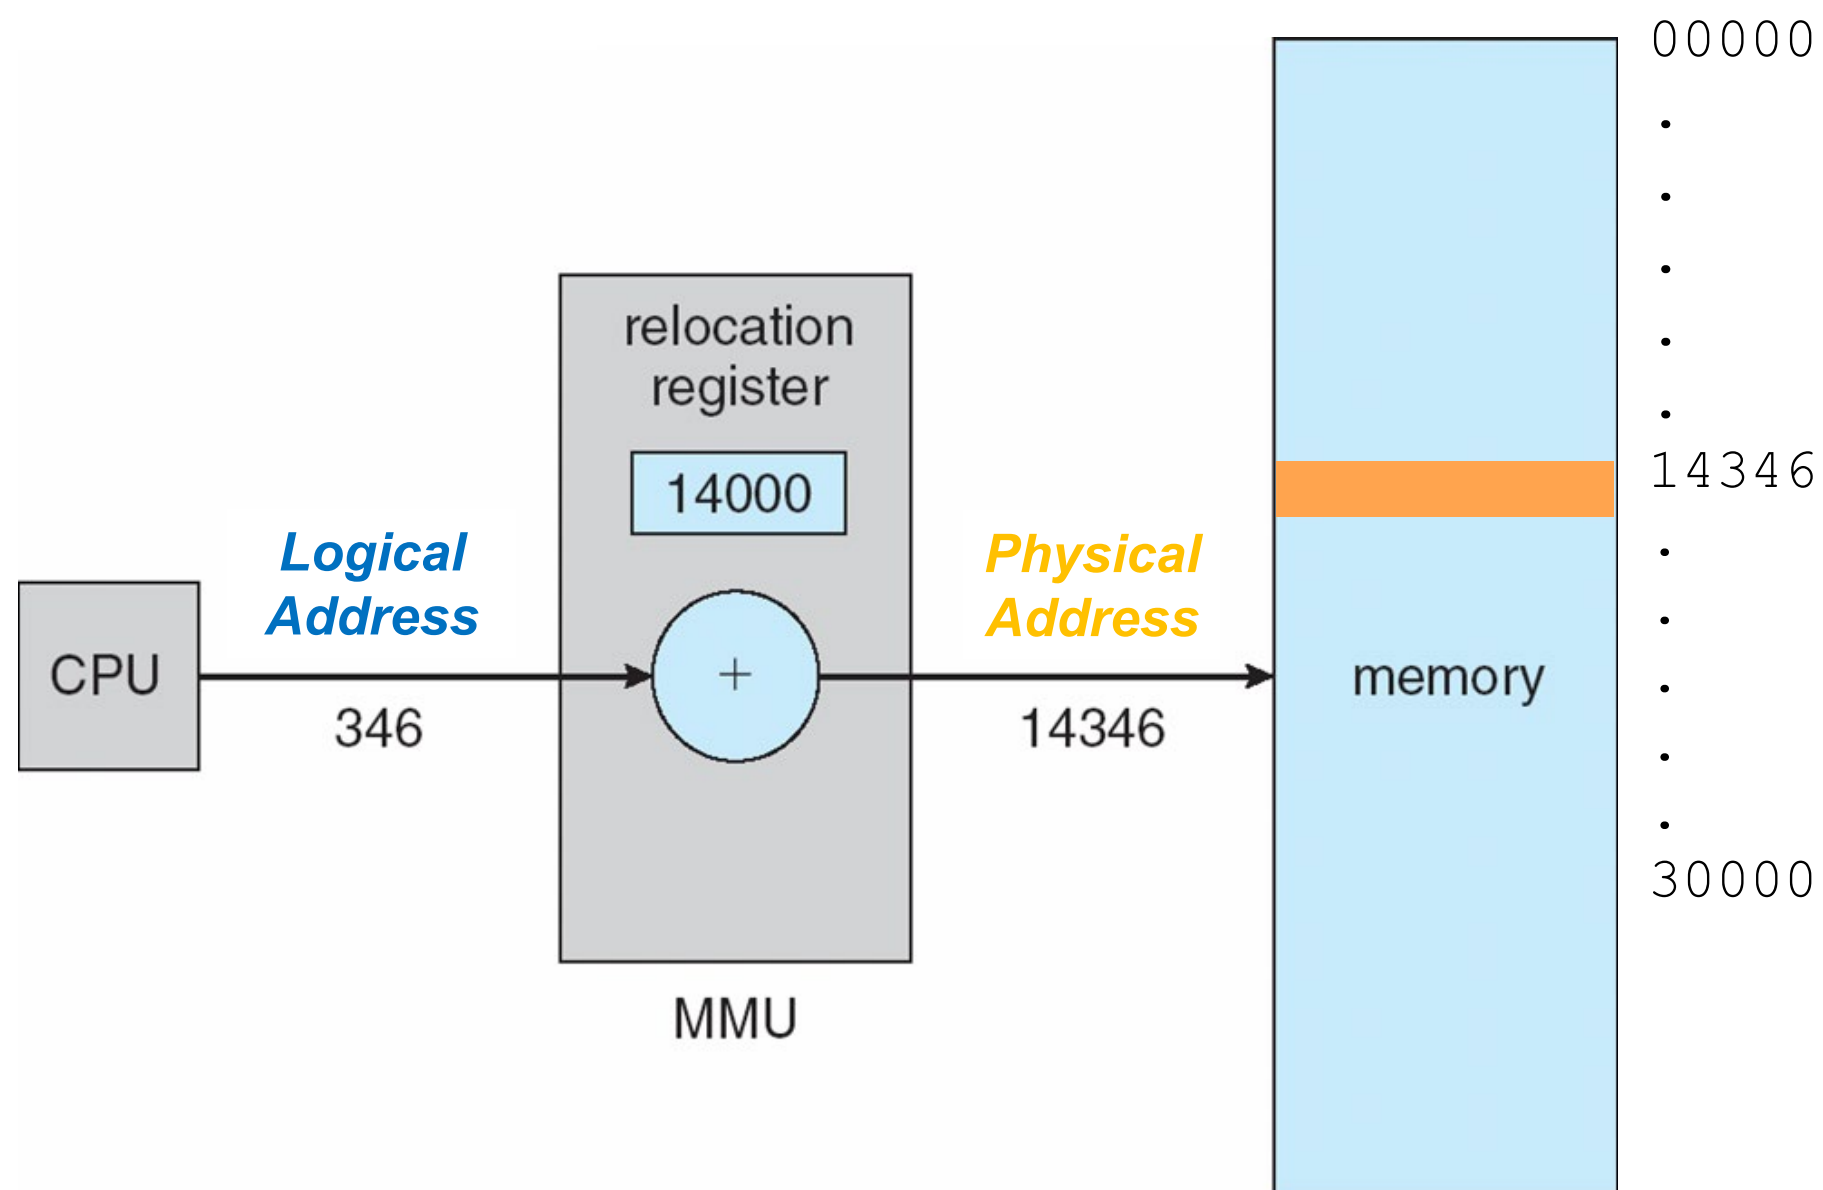
- Hardware device that at run time maps **logical / virtual addresses** to **physical address**.



- The user program deals with **logical addresses**; it never sees the real **physical addresses**.
  - Execution-time binding occurs when reference is made to location in memory.
  - **Logical address** bound to **physical addresses**







**Figure 8.4** Dynamic relocation using a relocation register.





## 2. SWAPPING


---

- A process can be **swapped** out of memory to a **backing store** (temporarily) and then brought back into memory for continued execution
  - Total physical memory space of all processes can exceed the real physical memory of the system.
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all processes; must provide direct access to these memory images.
- System maintains a **ready queue** of ready-to-run processes which are either in;
  - i) memory or
  - ii) have memory images on disk.





- Major part of swap time is **transfer time**; total transfer time is directly proportional to the amount of memory swapped.



**Figure 8.5** Swapping of two processes using a disk as a backing store.

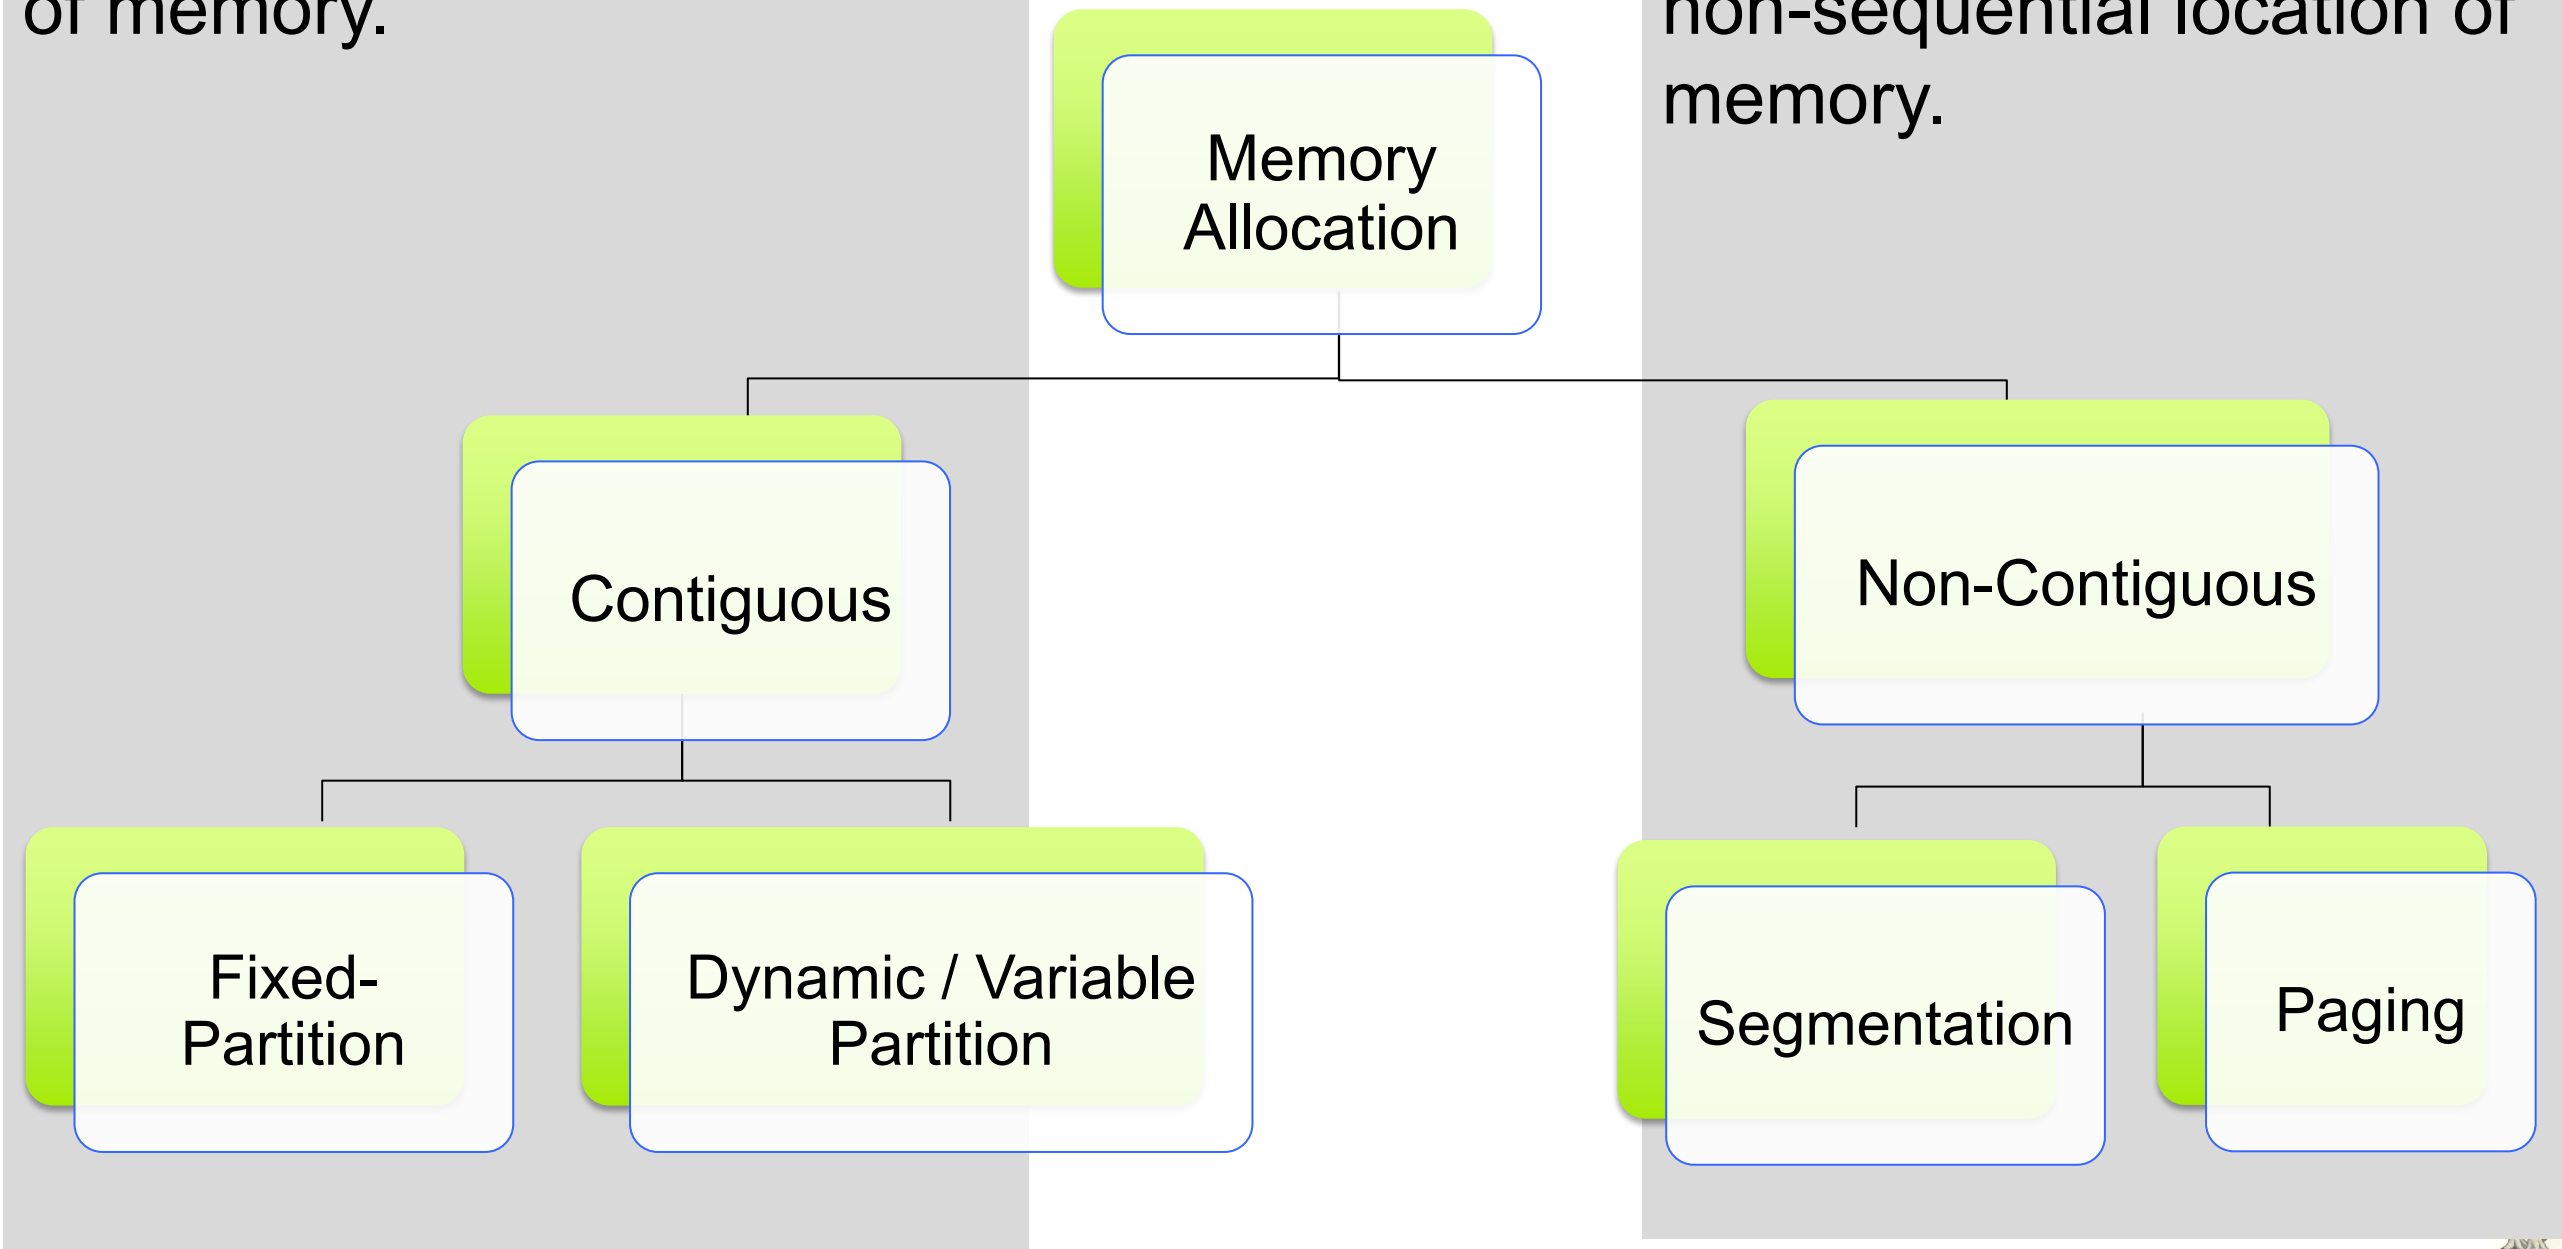




# 3. MEMORY ALLOCATION

Each process is contained in a single **contiguous** section of memory.

Each process is contained in a **non-contiguous** / non-sequential location of memory.





## 4. CONTIGUOUS MEMORY ALLOCATION

---

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one of early methods
  
- Main memory divided usually into two **partitions**:
  - One for **operating system**, usually held in low memory
  - **User processes** then held in high memory
  - Each process is contained in **single contiguous** section of memory





## (a) Fixed-sized Partition : Example

- Each partition may contain exactly one process.
- The degree of multiprogramming is limited by number of partitions
- But it requires
  - Protection of the job's memory space
  - Matching job size with partition size

A simplified fixed-partition memory table with the free partition shaded.

| Partition Size | Memory Address | Access | Partition Status |
|----------------|----------------|--------|------------------|
| 100K           | 200K           | Job 1  | Busy             |
| 25K            | 300K           | Job 4  | Busy             |
| 25K            | 325K           |        | Free             |
| 50K            | 350K           | Job 2  | Busy             |





# Example 1

| Partition Size | Memory Address | Access | Partition Status |
|----------------|----------------|--------|------------------|
| 100K           | 200K           | Job 1  | Busy             |
| 25K            | 300K           | Job 4  | Busy             |
| 25K            | 325K           |        | Free             |
| 50K            | 350K           | Job 2  | Busy             |

Job List:

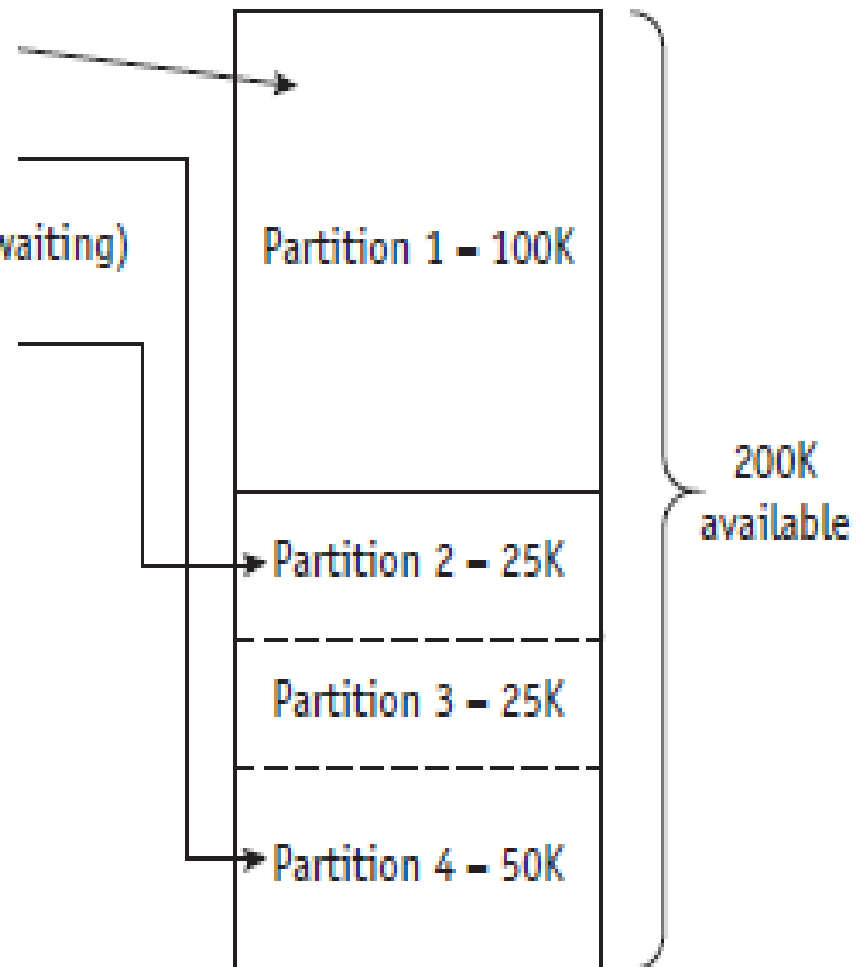
Job 1 – 30K

Job 2 – 50K

Job 3 – 30K (waiting)

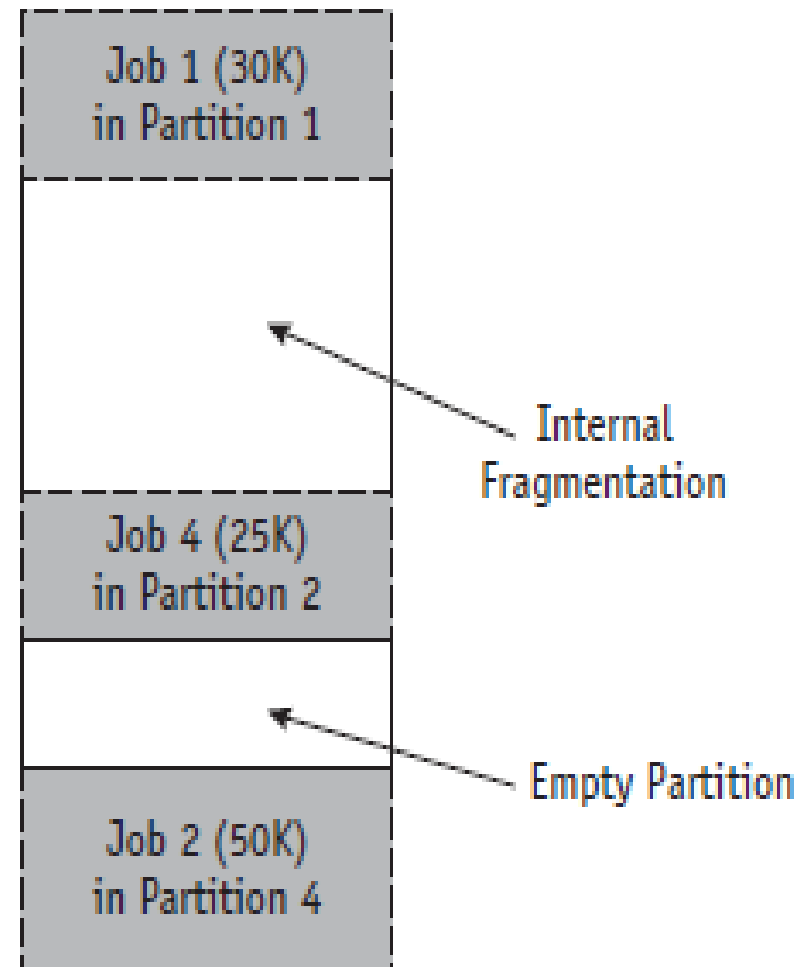
Job 4 – 25K

Main Memory



(a)

Main Memory



(b)

(figure 2.3)

*Main memory use during fixed partition allocation of Table 2.1. Job 3 must wait even though 70K of free space is available in Partition 1, where Job 1 only occupies 30K of the 100K available. The jobs are allocated space on the basis of “first available partition of required size.”*





- Disadvantages Fixed Partitions
  - Requires contiguous loading of entire program.
  - Job allocation method:
    - *First available partition with required size.*
  - Arbitrary partition size leads to undesired results:
    - Partition too small
      - Large jobs have **longer turnaround time.**
    - Partition too large
      - Memory waste: **internal fragmentation.**








# Exercise 1

- The new jobs in the order of ...
  - Job *P* (20K) – Job *Q* (15K) – Job *R* (10K) – Job *S* (25K) need to be allocated in a **fixed partition** memory as shown in the figure.
  - Label all the jobs in which partition each will be allocated.

|                    |    |     |
|--------------------|----|-----|
| <i>Partition 0</i> | OS |     |
| <i>Partition 1</i> |    | 20K |
| <i>Partition 2</i> |    | 5K  |
| <i>Partition 3</i> |    | 10K |
| <i>Partition 4</i> |    | 20K |
| <i>Partition 5</i> |    | 5K  |
| <i>Partition 6</i> |    | 10K |
| <i>Partition 7</i> |    | 30K |





# Solution

Fixed partition: Job P (20K)  
Job Q (15K)  
Job R (10K)  
Job S (25K)

|             |    |     |
|-------------|----|-----|
| Partition 0 | OS |     |
| Partition 1 |    | 20K |
| Partition 2 |    | 5K  |
| Partition 3 |    | 10K |
| Partition 4 |    | 20K |
| Partition 5 |    | 5K  |
| Partition 6 |    | 10K |
| Partition 7 |    | 30K |

|             |       |     |
|-------------|-------|-----|
| Partition 0 | OS    |     |
| Partition 1 | Job P | 20K |
| Partition 2 |       | 5K  |
| Partition 3 | Job R | 10K |
| Partition 4 | Job Q | 20K |
| Partition 5 |       | 5K  |
| Partition 6 |       | 10K |
| Partition 7 | Job S | 30K |

Internal fragmentation (5K)

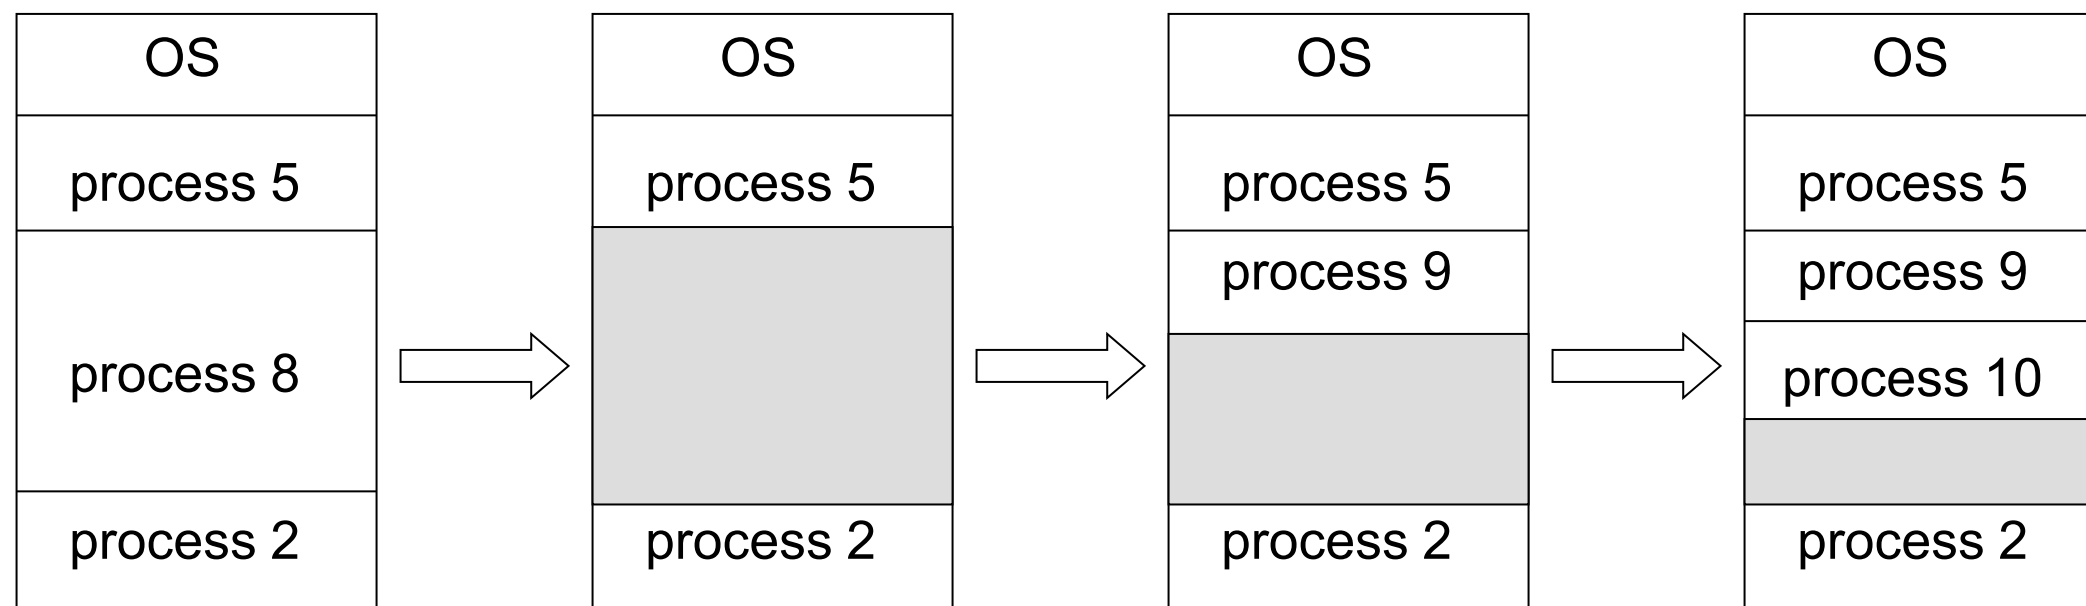
Internal fragmentation (5K)





## (b) Dynamic/Variable Partition

- **Dynamic/Variable-partition** sizes for efficiency (sized to a given process' needs)
  - **Hole** – block of available memory; holes of various size are scattered throughout memory
  - When a process arrives, it is allocated memory from a **hole large enough** to accommodate it
  - Process exiting **fre**es its partition, adjacent free partitions will be **combined**
  - Operating system maintains information about:
    - a) allocated partitions
    - b) free partitions (hole)



 *Holes / External fragmentation*





# Example 2

**Variable/Dynamic partition:** Five snapshots (a – e) of main memory as 8 jobs submitted and allocated space on the basis “***First-Come, First-Serve***”.

Init. jobs:

*A* (10K)

*B* (15K)

*C* (20K)

*D* (50K)

New jobs:


*E* (5K)

*F* (30K)

New jobs:


*G* (10K)

*H* (30K)




(a)

Initial job entry memory allocation.



(b)

After Job *A* and *D* have finished.



(c)

After Job *E* and *F* have entered.





**Variable/Dynamic partition:** Five snapshots (a – e) of main memory as 8 jobs submitted and allocated space on the basis “***First-Come, First-Serve***”.

Init. jobs:

*A* (10K)

*B* (15K)

*C* (20K)

*D* (50K)

New jobs:

*E* (5K)

*F* (30K)

New jobs:


*G* (10K)

*H* (30K)



(d)

After Job *C* has finished.



(e)

After Job *G* has entered.





# Solution

**Variable/Dynamic partition:** Five snapshots (a – e) of main memory as 8 jobs submitted and allocated space on the basis “*First-Come, First-Serve*”.

Init. jobs:


*A* (10K)  
*B* (15K)  
*C* (20K)  
*D* (50K)

New jobs:

*E* (5K)  
*F* (30K)


New jobs:

*G* (10K)  
*H* (30K)




(a)

Initial job entry memory allocation.



(b)

After Job *A* and *D* have finished.



(c)

After Job *E* and *F* have entered.





Init. jobs:


*A* (10K)  
*B* (15K)  
*C* (20K)  
*D* (50K)

New jobs:

*E* (5K)  
*F* (30K)


New jobs:

*G* (10K)  
*H* (30K)



(d)

After Job *C* has  
finished.



(e)

After Job *G* has  
entered.

**Job *H*** has to wait even  
though there's enough  
free memory in  
between partitions to  
accommodate it.





# Exercise 2

Figure shows a few jobs have been allocated in a **dynamic partition** memory at time  $t_0$ . Suppose that job  $P$  and  $Q$  finished at  $t_1$  and new jobs arrived at  $t_3$ . Allocate all the new jobs in the memory.

## New jobs:

$T (30K)$

$U (35K)$

$V (15K)$

|         |     |
|---------|-----|
| OS      |     |
| Job $P$ | 20K |
|         | 5K  |
| Job $R$ | 10K |
| Job $Q$ | 15K |
|         | 20K |
| Job $S$ | 25K |
|         | 5K  |







# Solution

$t_0$ : Initial jobs in memory.

$t_1$ : After Job  $P$  and  $Q$  have finished.

$t_3$ : After Job  $T$  and  $V$  have entered.

Job  $U$  need to wait.

## New jobs:

$T$  (30K)  
 $U$  (35K)  
 $V$  (15K)

| OS      |     |
|---------|-----|
| Job $P$ | 20K |
|         | 5K  |
| Job $R$ | 10K |
| Job $Q$ | 15K |
|         | 20K |
| Job $S$ | 25K |
|         | 5K  |

| OS      |     |
|---------|-----|
|         | 25K |
| Job $R$ | 10K |
|         | 35K |
| Job $S$ | 25K |
|         | 5K  |

| OS      |     |
|---------|-----|
| Job $V$ | 15K |
|         | 10K |
| Job $R$ | 10K |
| Job $T$ | 30K |
|         | 5K  |
| Job $S$ | 25K |
|         | 5K  |





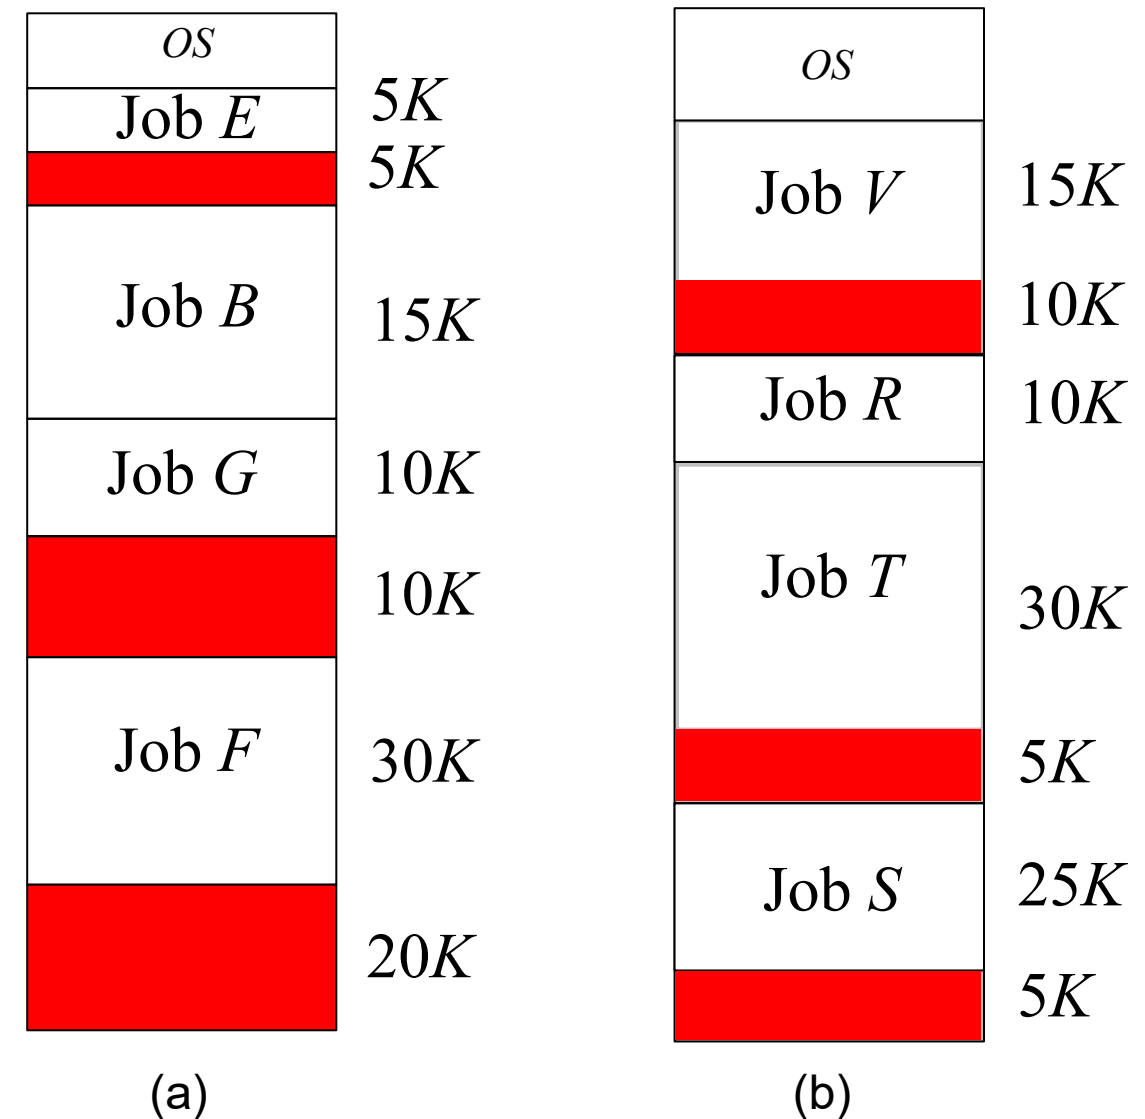
- Based on previous examples, job *H* and *U* respectively have to wait until any job ends with enough size.

- Disadvantages:**

Produce block of available memory (**holes**) with different sizes that scattered out the memory.

- Solution:**

**Relocatable Dynamic Partition**



Holes or external fragmentation examples before





# Fragmentation

- There are two types of fragmentation:

## Internal fragmentation

- allocated memory may be slightly larger than requested memory;
- the difference of memory internal to a partition that not being used;
- Results from **Fixed partition allocation**.

## External fragmentation


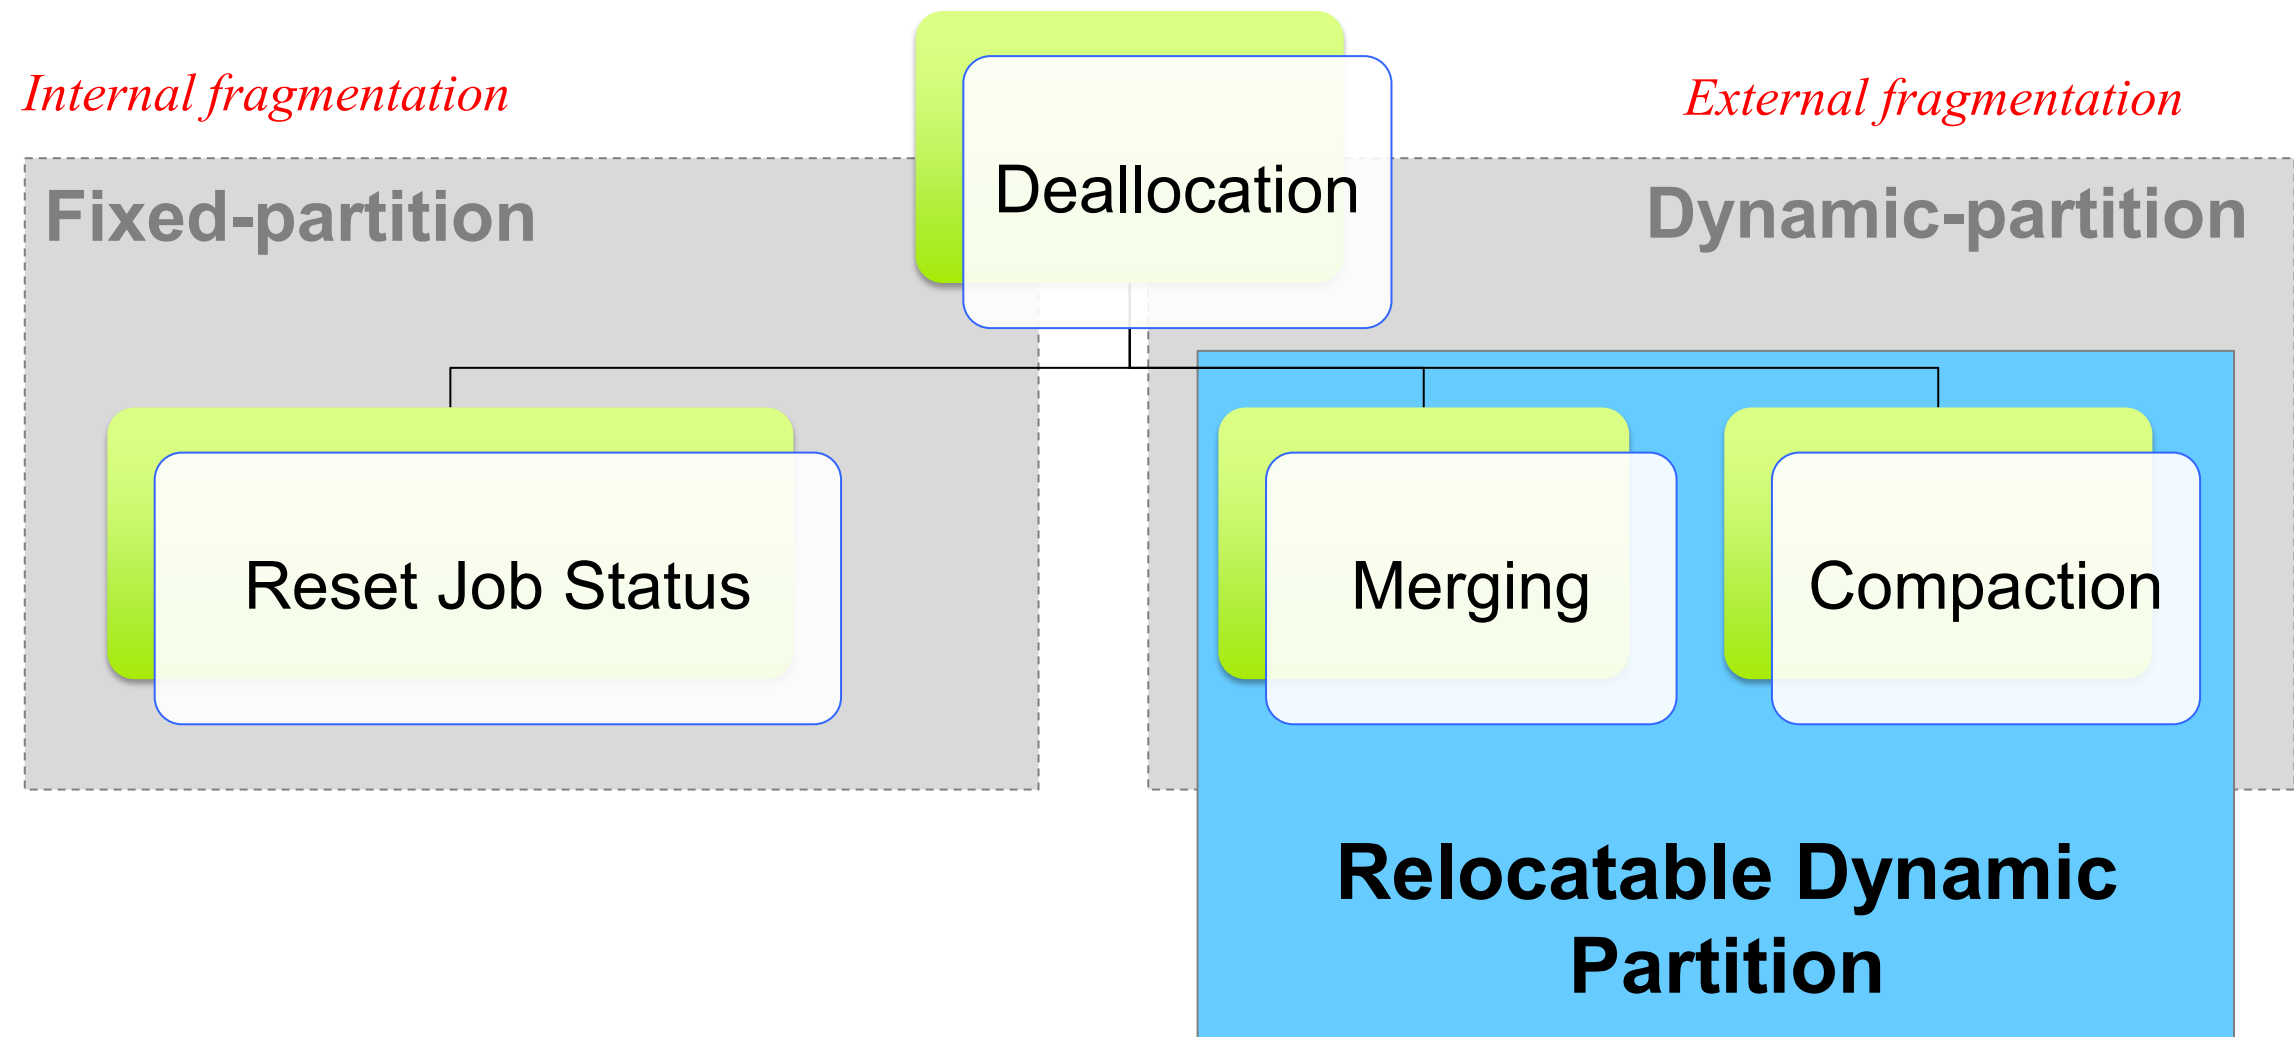
- total memory space exists to satisfy a request, but it is not contiguous.
- Results from **Dynamic/Variable partition allocation**.





# Deallocation

- Solution to reduce **fragmentation**  
→ **Deallocation** : freeing allocated memory spaces.





- For **fixed-partition** system:
  - Straightforward process
  - Memory manager **resets** the status of job's memory block to **free** upon job completion
  - Example
    - Binary values with
      - **0** indicating **FREE**, and **1** indicating **BUSY**

A simplified fixed-partition memory table with the free partition shaded.

| Partition Size | Memory Address | Access | Partition Status |
|----------------|----------------|--------|------------------|
| 100K           | 200K           | Job 1  | Busy             |
| 25K            | 300K           | Job 4  | Busy             |
| 25K            | 325K           |        | Free             |
| 50K            | 350K           | Job 2  | Busy             |





- For **dynamic-partition** system:
  - Algorithm tries to merge free areas of memory
  - More complex
- Reduce external fragmentation by **merging and compaction**
  - **Compaction:**
    - reclaiming fragmented sections of memory space.
    - shuffle memory contents to place all free memory together in one large block.
    - every program in memory must be relocated.
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time
  - It is called **Relocatable Dynamic Partition**

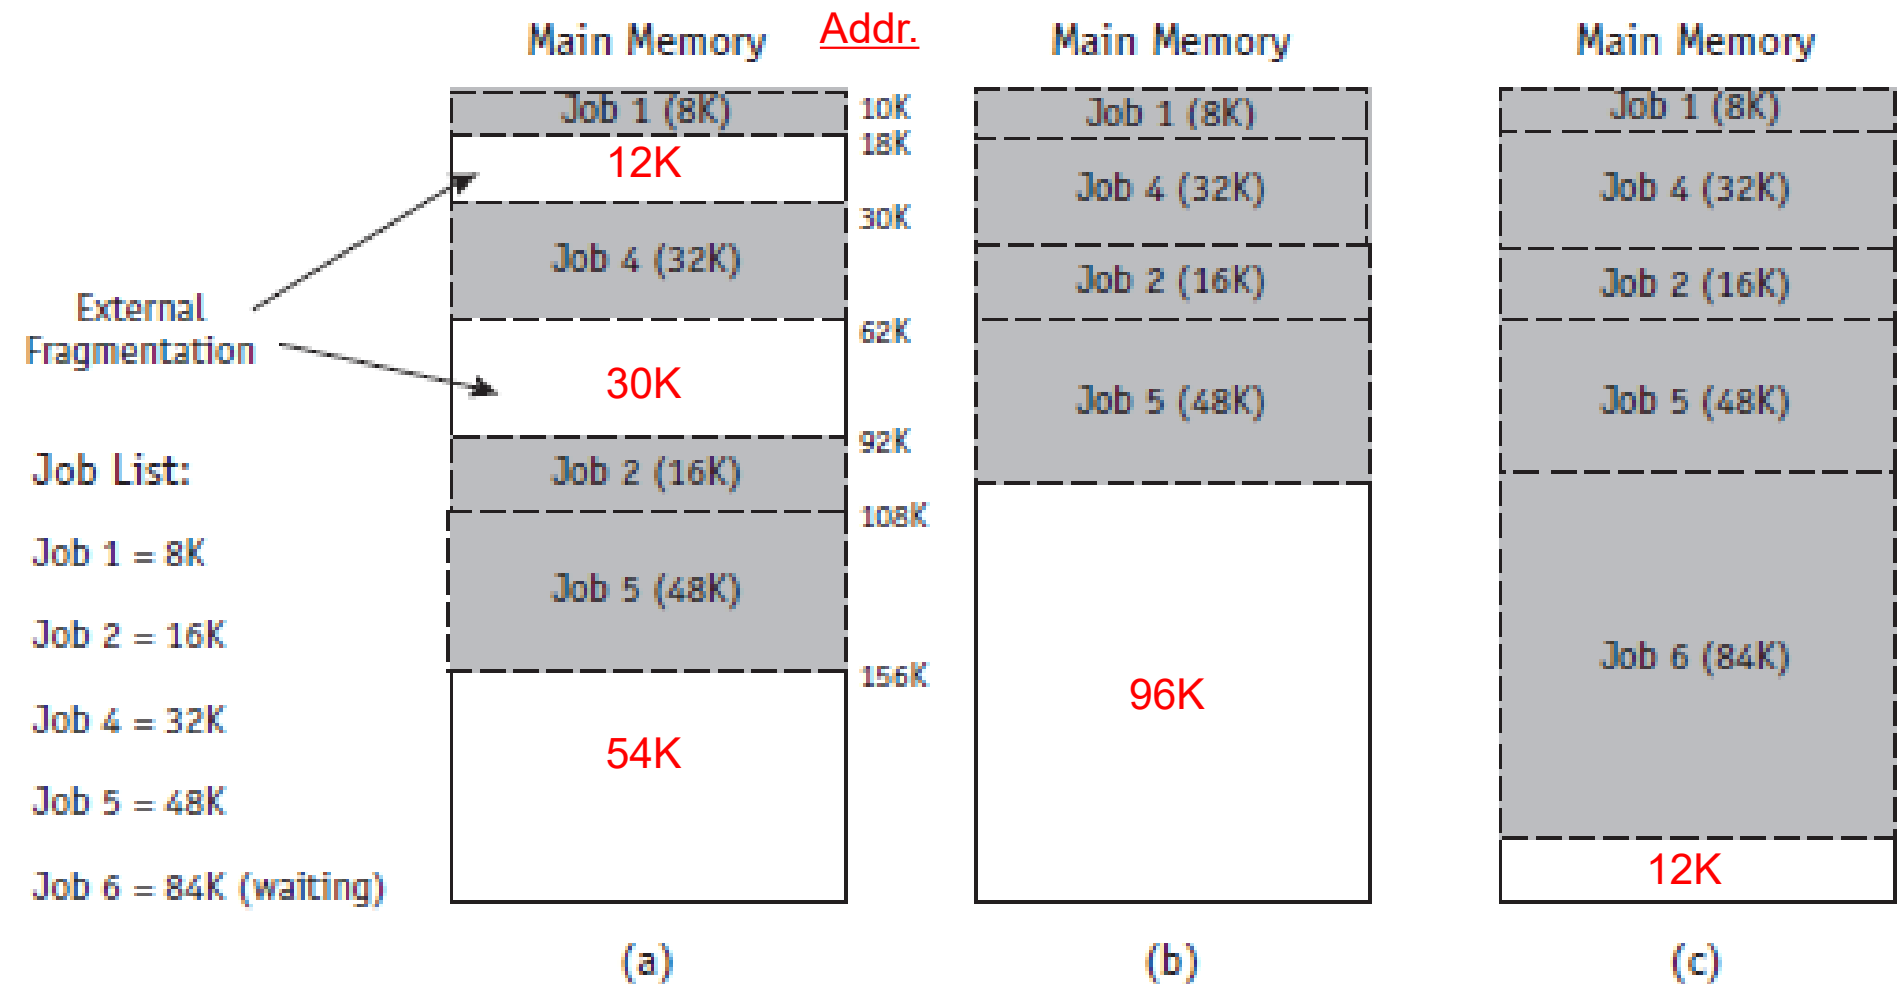




# Example 3

## Solution to External Fragmentation – Compaction

**Merge or combine  
all scattered holes at  
one location.**



(figure 2.9)

Three snapshots of memory before and after compaction with the operating system occupying the first 10K of memory. When Job 6 arrives requiring 84K, the initial memory layout in (a) shows external fragmentation totaling 96K of space. Immediately after compaction (b), external fragmentation has been eliminated, making room for Job 6 which, after loading, is shown in (c).





# Example 4

**Relocatable Dynamic partition:** Three snapshots (a – c) of main memory (dynamic partition system) before and after **compaction**.

Jobs:

*B* (15K)

*E* (5K)

*F* (30K)

*H* (30K) - waiting

|              |     |
|--------------|-----|
| <i>OS</i>    |     |
| Job <i>E</i> | 5K  |
|              | 5K  |
| Job <i>B</i> | 15K |
|              | 20K |
| Job <i>F</i> | 30K |
|              | 20K |

(a)

External fragmentation total 45K.







# Solution


## Jobs:

*B* (15K)

*E* (5K)


*F* (30K)

*H* (30K) - waiting




(a)

External fragmentation total 45K.




(b)

Compaction and merge to make a large empty space



(c)

Job *H* entered the memory.

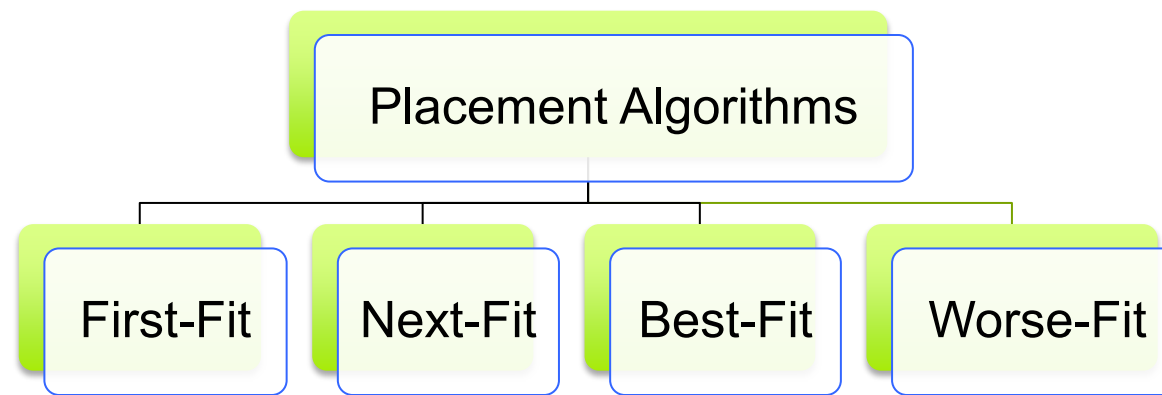




# Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

- Use the **Placement Algorithms** for selecting among free blocks of main memory.

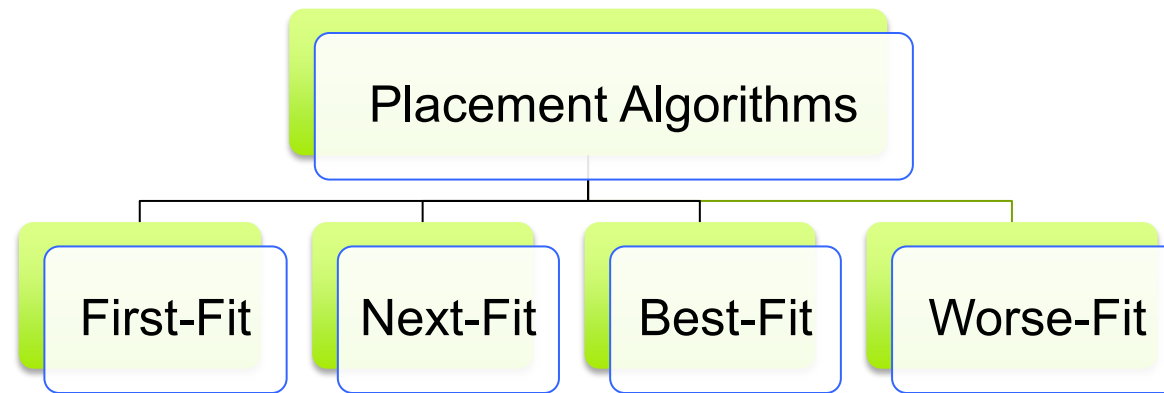




# Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

- Use the **Placement Algorithms** for selecting among free blocks of main memory.



- **First-fit:**
  - Allocate the **first hole** that is big enough.
  - Search start at the beginning of memory.
- **Next-fit:**
  - Similar to first-fit, except the search starts at the **last hole** allocated.

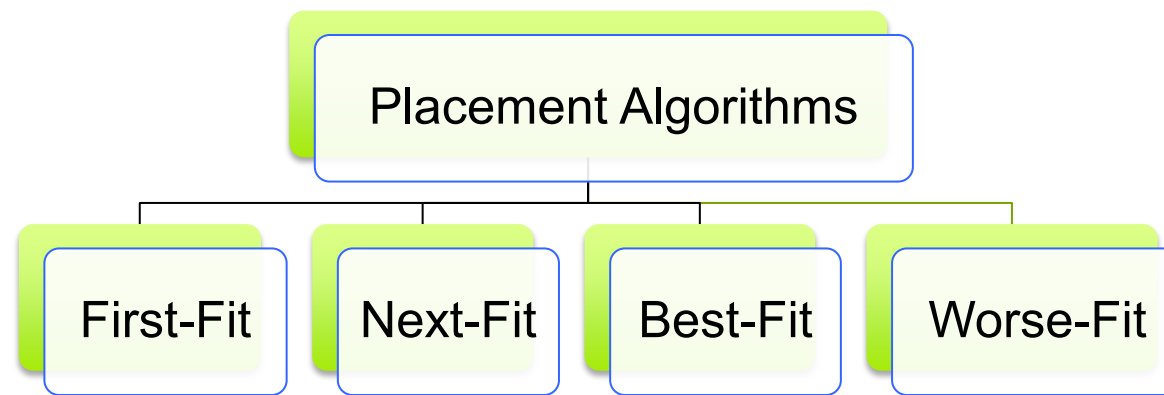




# Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

- Use the **Placement Algorithms** for selecting among free blocks of main memory.



- **Best-fit:**
  - Allocate the ***smallest hole*** that is big enough;
  - must search entire list, unless ordered by size
  - Produces the smallest leftover hole
- **Worst-fit:**
  - Allocate the ***largest hole***; must also search entire list
  - Produces the largest leftover hole





# First-Fit Allocation

Example 5 : Assume fixed partition are used

| <i>Main Memory</i> |    |     |
|--------------------|----|-----|
| <i>Partition 0</i> | OS |     |
| <i>Partition 1</i> |    | 30K |
| <i>Partition 2</i> |    | 15K |
| <i>Partition 3</i> |    | 50K |
| <i>Partition 4</i> |    | 20K |

Init. jobs:

*A* (10K)

*B* (20K)

*C* (30K)


*D* (10K)





## Solution 5 :

**First-fit allocation:**  
(Assume fixed partition are used)



### Init. jobs:

A (10K)

B (20K)

C (30K)

D (10K)

- Job C has to wait until a large block become available, even though there is 85K unused memory space.

Unused partition (20K)

Internal Fragmentation (55K)





# Example 6: Best-Fit Allocation

Example 6 : Assume fixed partition are used

| <i>Main Memory</i> |           |            |
|--------------------|-----------|------------|
| <i>Partition 0</i> | <i>OS</i> |            |
| <i>Partition 1</i> |           | <i>30K</i> |
| <i>Partition 2</i> |           | <i>15K</i> |
| <i>Partition 3</i> |           | <i>50K</i> |
| <i>Partition 4</i> |           | <i>20K</i> |

Init. jobs:

*A (10K)*

*B (20K)*

*C (30K)*

*D (10K)*






## Solution 6 :

### Best-fit allocation:

(Assume fixed partition are used)



#### Init. jobs:

*A (10K)*

*B (20K)*

*C (30K)*

*D (10K)*

- ✓ All jobs can be allocated in the memory.
- ✓ Use the memory efficiently but **slower** implementation.







# Example 7: Next-Fit Allocation

Example 7 : Assume fixed partition are used

| <i>Main Memory</i> |           |            |
|--------------------|-----------|------------|
| <i>Partition 0</i> | <i>OS</i> |            |
| <i>Partition 1</i> |           | <i>30K</i> |
| <i>Partition 2</i> |           | <i>15K</i> |
| <i>Partition 3</i> |           | <i>50K</i> |
| <i>Partition 4</i> |           | <i>20K</i> |

Init. jobs:

*A (10K)*

*B (20K)*

*C (30K)*

*D (10K)*






## Solution 7 :

### Next-fit allocation:

(Assume fixed partition are used)



#### Init. jobs:

*A (10K)*

*B (20K)*

*C (30K)*

*D (10K)*

➤ Job *C* has to wait until a any job exit and has enough 30K memory space.

✓ This scheme also use the memory efficiently but **slower** implementation.

 *Unused partition (15K)*

 *Internal Fragmentation (60K)*





# Example 8: Worst-Fit Allocation

Example 8 : Assume fixed partition are used

| <i>Main Memory</i> |           |            |
|--------------------|-----------|------------|
| <i>Partition 0</i> | <i>OS</i> |            |
| <i>Partition 1</i> |           | <i>30K</i> |
| <i>Partition 2</i> |           | <i>15K</i> |
| <i>Partition 3</i> |           | <i>50K</i> |
| <i>Partition 4</i> |           | <i>20K</i> |

Init. jobs:

*A (10K)*

*B (20K)*

*C (30K)*


*D (10K)*





Solution 8 :

**Worst-fit allocation:**  
(Assume fixed partition are used)



Init. jobs:

A (10K)

B (20K)

C (30K)

D (10K)

➤ Job C has to wait until a large block become available

Unused partition (15K)

Internal Fragmentation (60K)





# Best-Fit vs First-Fit Allocation

- **First-fit memory allocation**
  - Advantage: faster in making allocation
  - Disadvantage: **leads to memory waste**
- **Best-fit memory allocation**
  - Advantage: makes the best use of memory space
  - Disadvantage: **slower** in making allocation


First-fit and Best-fit better than Worst-fit in terms of speed and storage utilization.





## Example 9

Given a memory allocation as specified below.



Assume partition requests are in the following order:

$$P_1 - 20K, P_2 - 10K, P_3 - 5K$$


Determine which partition each request will be allocated if the following strategy is used:

(i) First-fit (ii) Next-fit (iii) Best-fit (iv) Worst-fit







## Solution 9:




First-fit




Best-fit



Next-fit



Worst-fit





# Summary

---

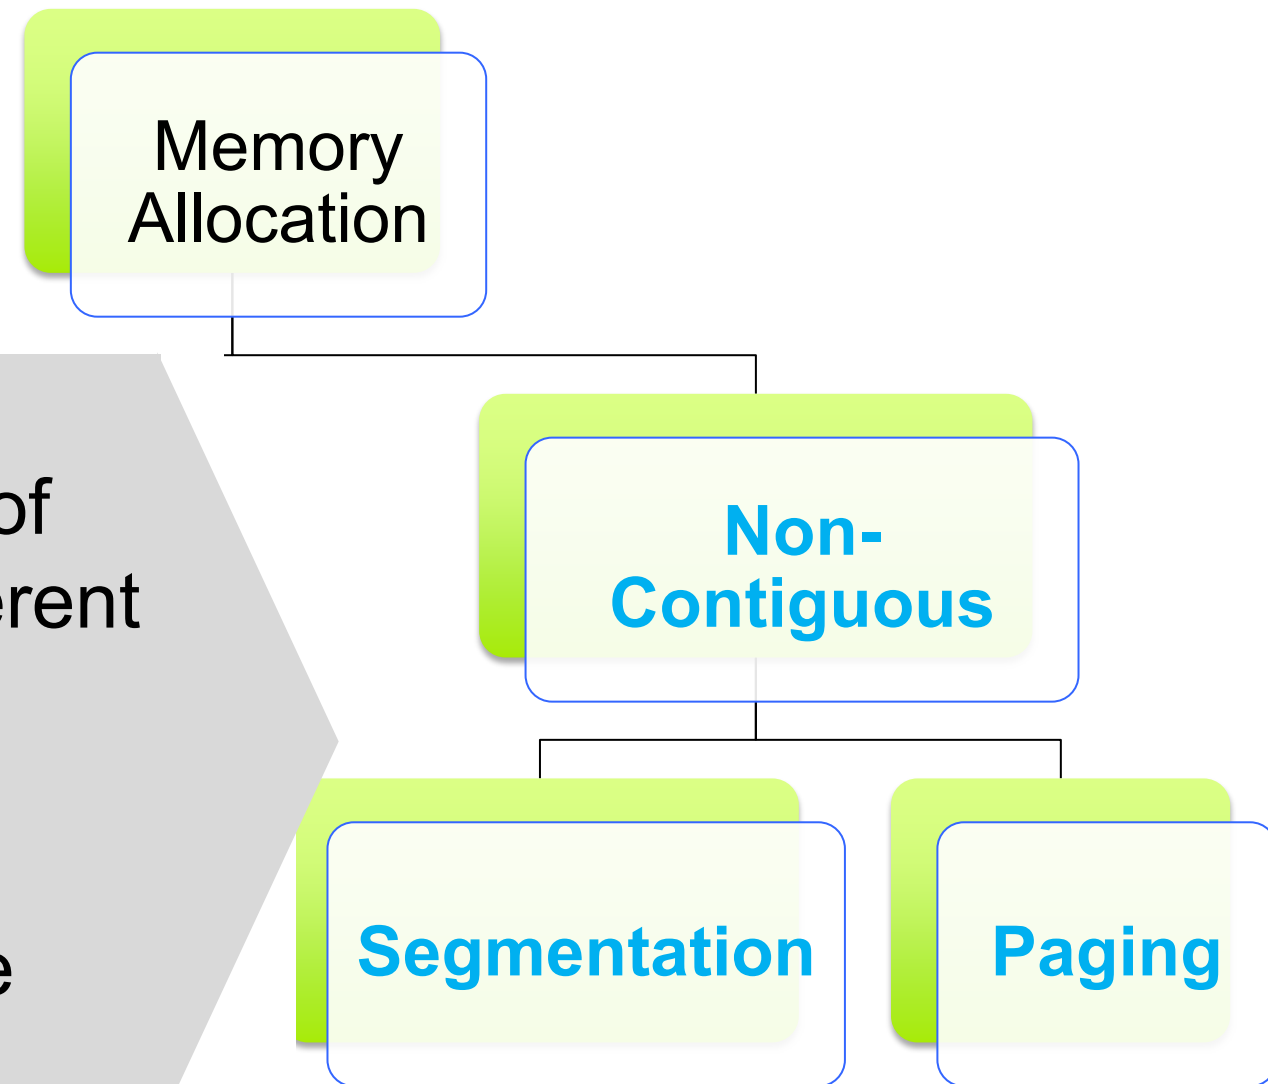
- Early memory management techniques
  - Fixed partitions, dynamic partitions, and relocatable dynamic partitions
- Common requirements of four memory management techniques
  - Entire program loaded into memory
  - Contiguous storage
  - Memory residency until job completed
- Each places severe restrictions on job size
- Sufficient for first three generations of computers
- New modern memory management trends
  - Common characteristics of memory schemes
    - ▶ Programs are not stored in contiguous memory
    - ▶ Not all segments reside in memory during job execution






# 5. NON CONTIGUOUS MEMORY ALLOCATION

- Partition a program into several small units, each of which can reside in a different parts of the memory.
- Need hardware support.
- Various methods to do the partitions:





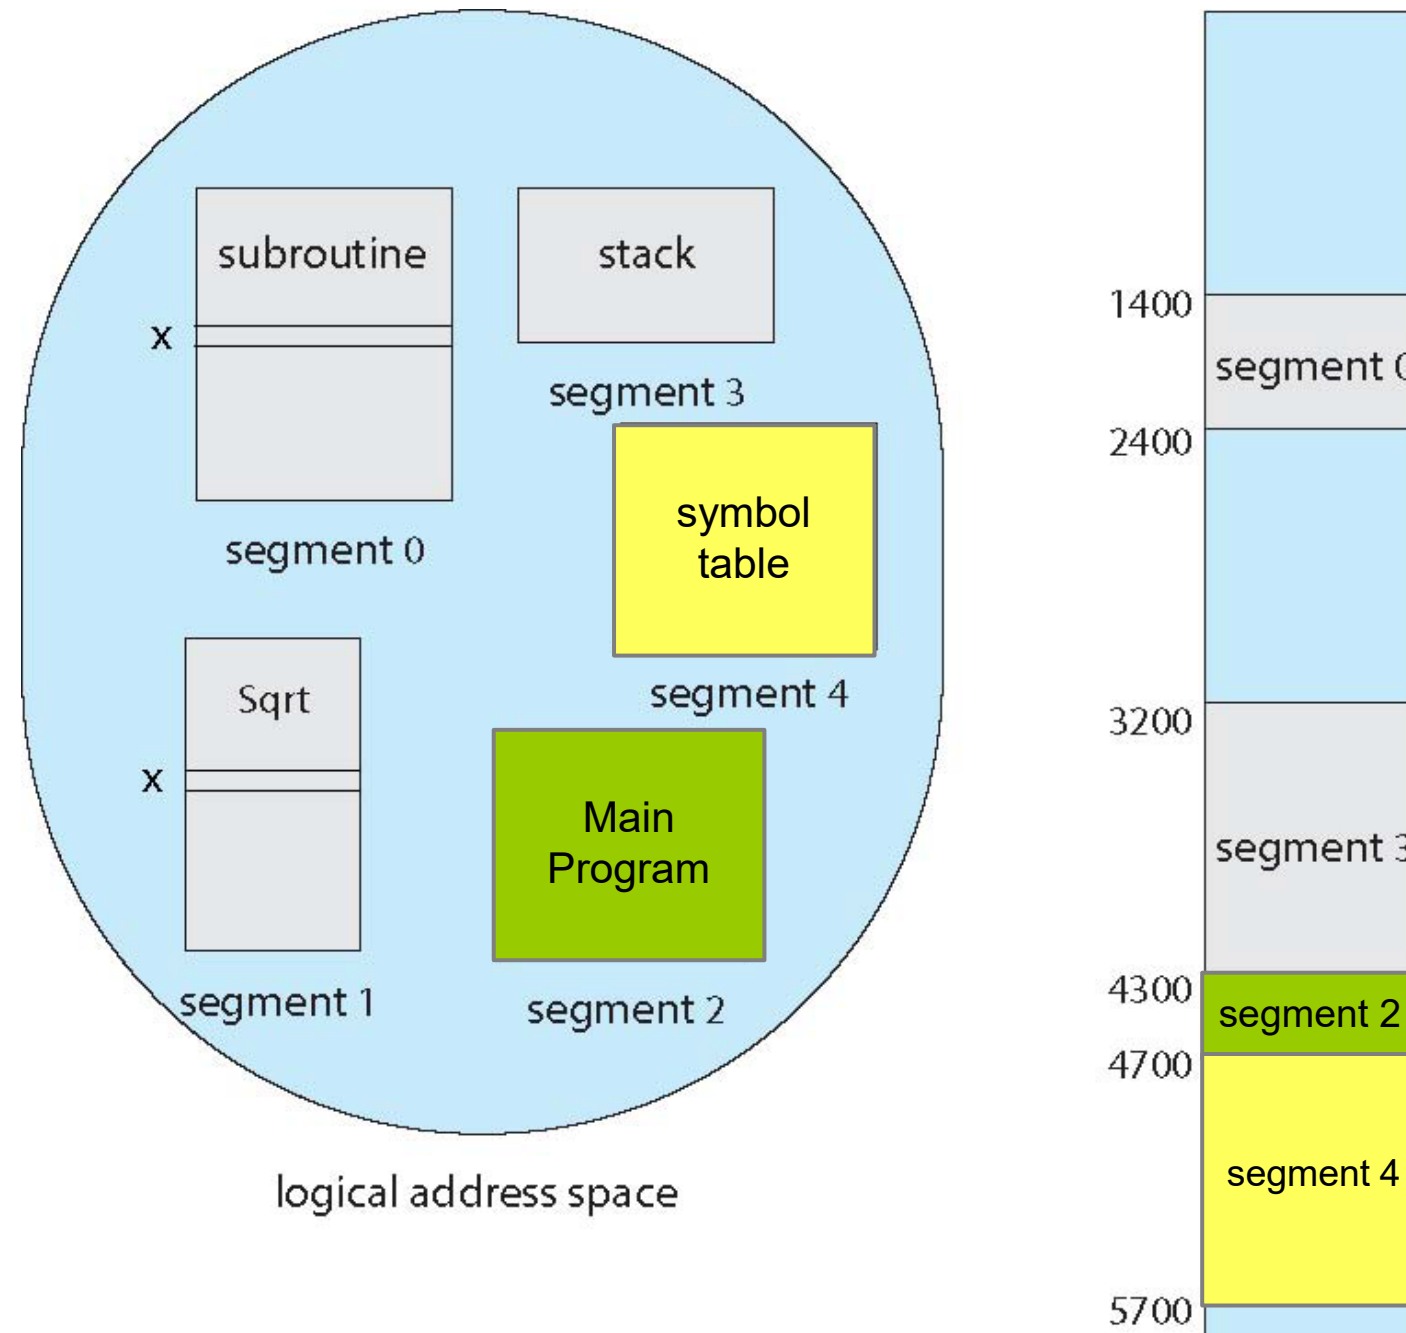
# (a) Segmentation



- Memory-management scheme that supports user view of memory.
- A program is a collection of **segments**.
- A segment is a **logical unit** such as:

|                      |                          |
|----------------------|--------------------------|
| <i>main program,</i> | <i>global variables,</i> |
| <i>procedure,</i>    | <i>common block,</i>     |
| <i>function,</i>     | <i>stack,</i>            |
| <i>method,</i>       | <i>local variables,</i>  |
| <i>object</i>        | <i>symbol table,</i>     |





**Figure 8.2** Logical and Physical Memory



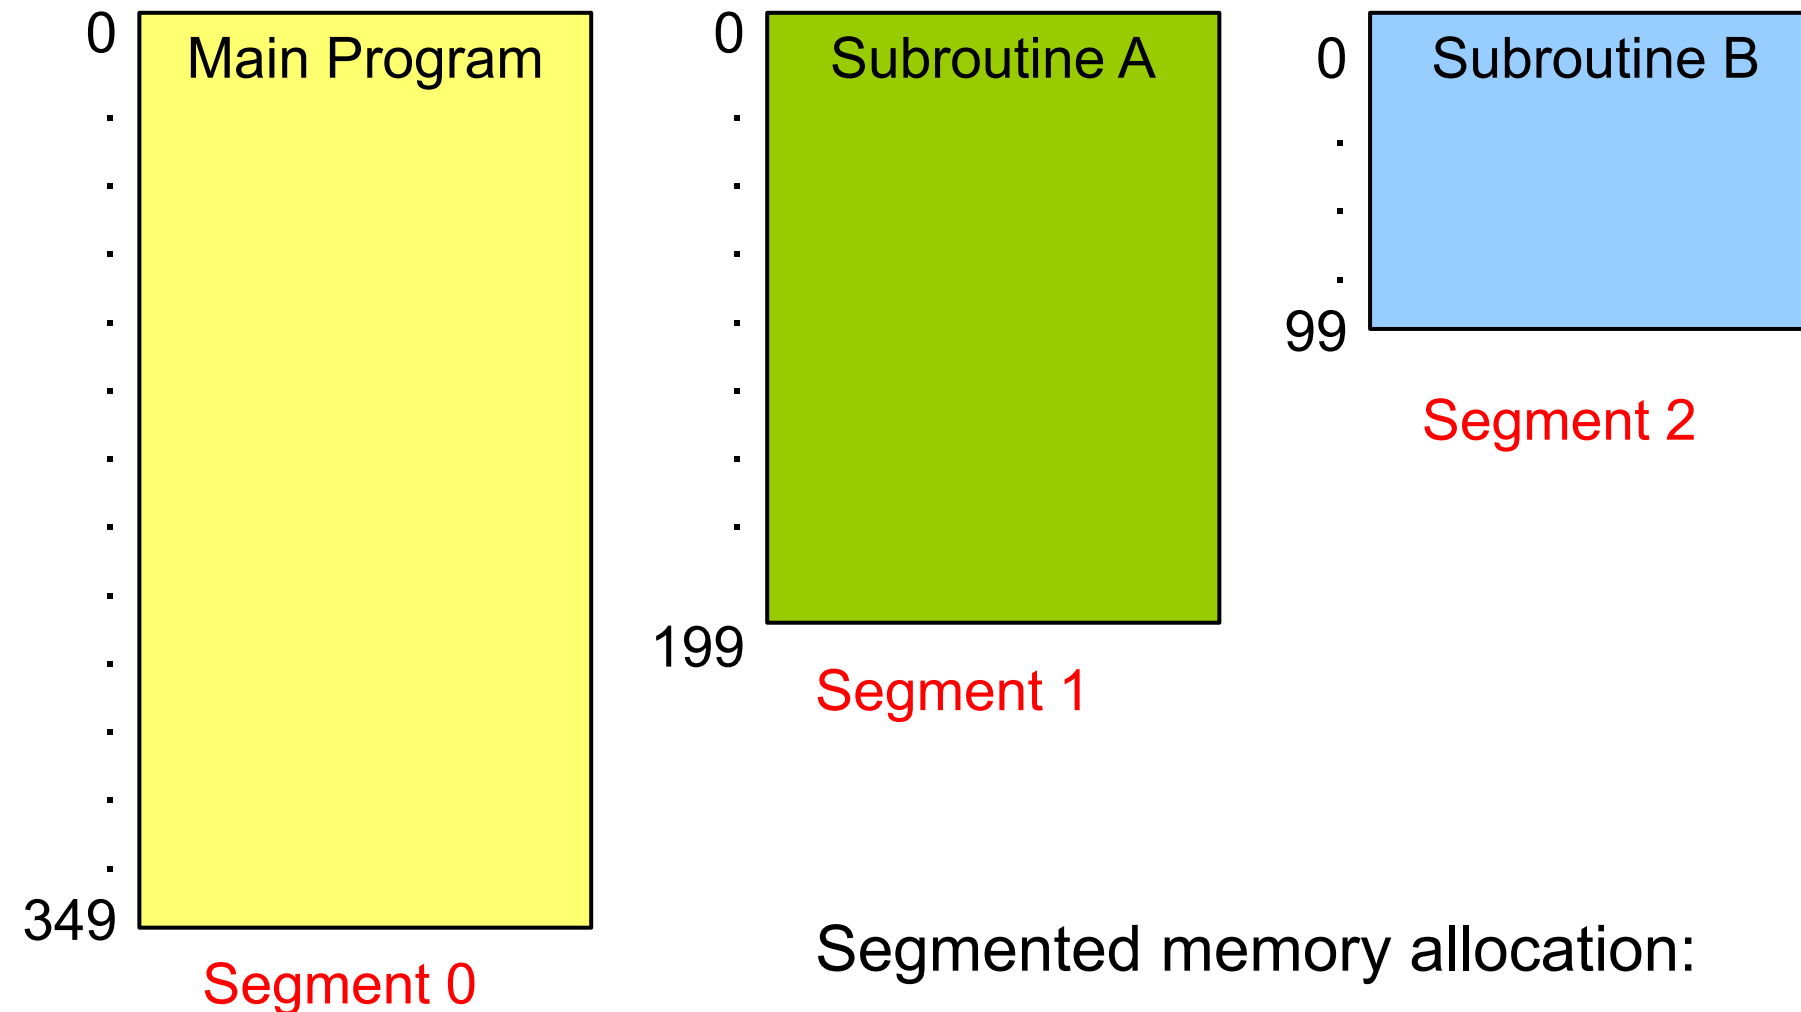


# Segmented Memory Allocation

---

- Each job is divided into several segments.
  - Segments are different sizes.
  - One for each module containing related functions.
- Main memory is allocated dynamically.
- Program's structural modules determine segments.
  - Each segment is numbered when compiled/assembled.
  - Need **Segment Table**.

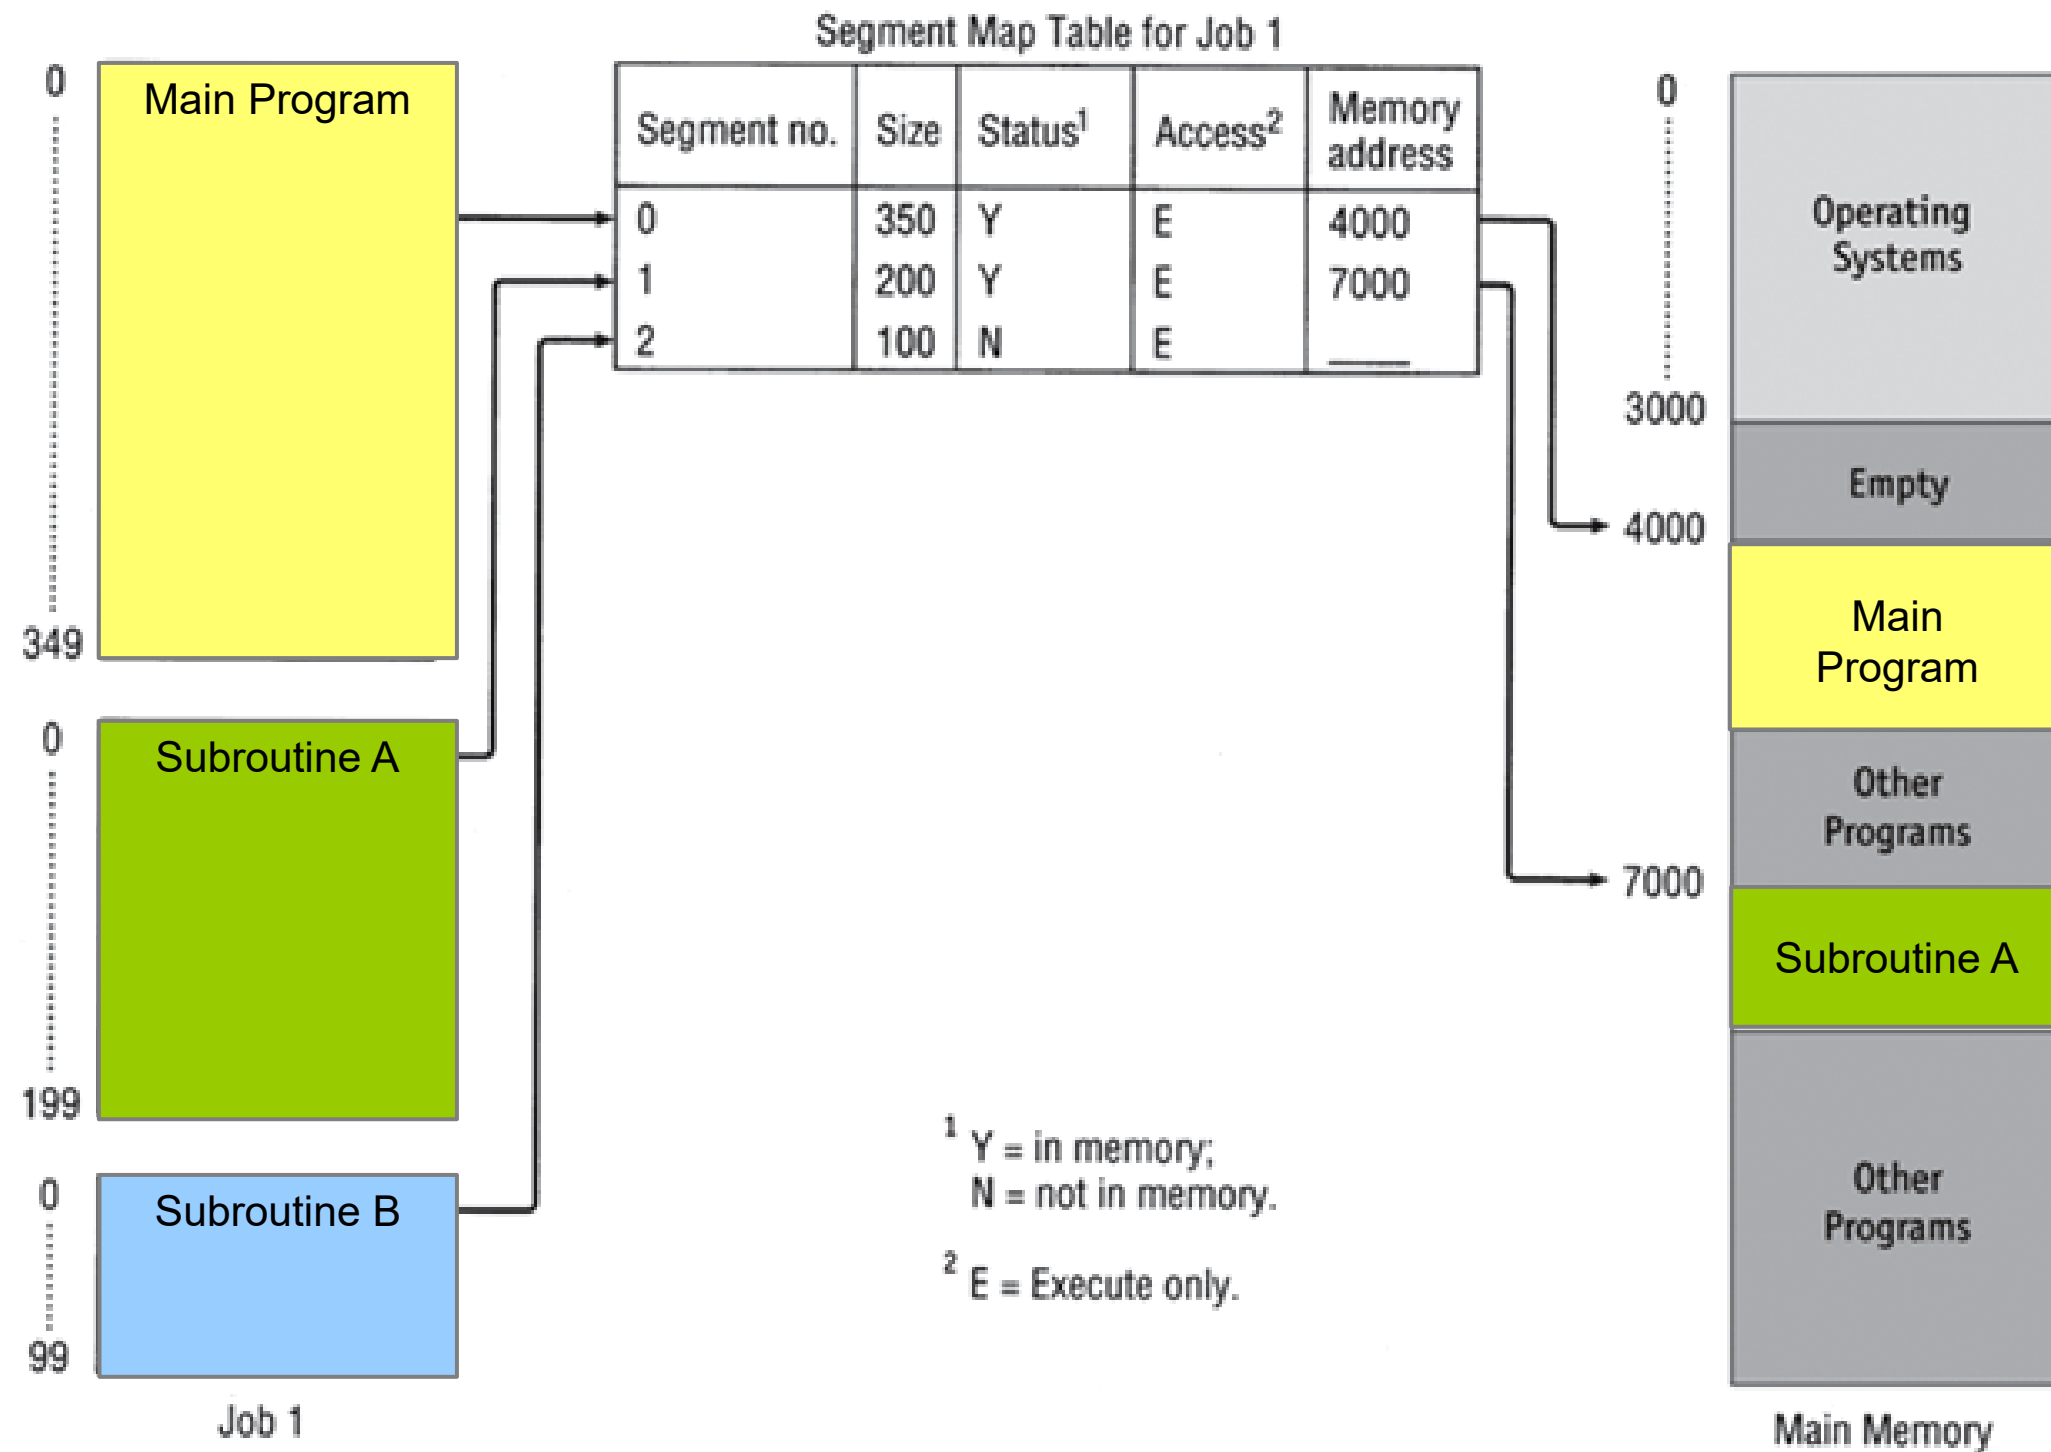




Segmented memory allocation:

- The job includes a Main Program, Subroutine *A*, and Subroutine *B*.
- It is one job divided into three segments.





(figure 3.14) The Segment Map Table tracks each segment for Job 1.





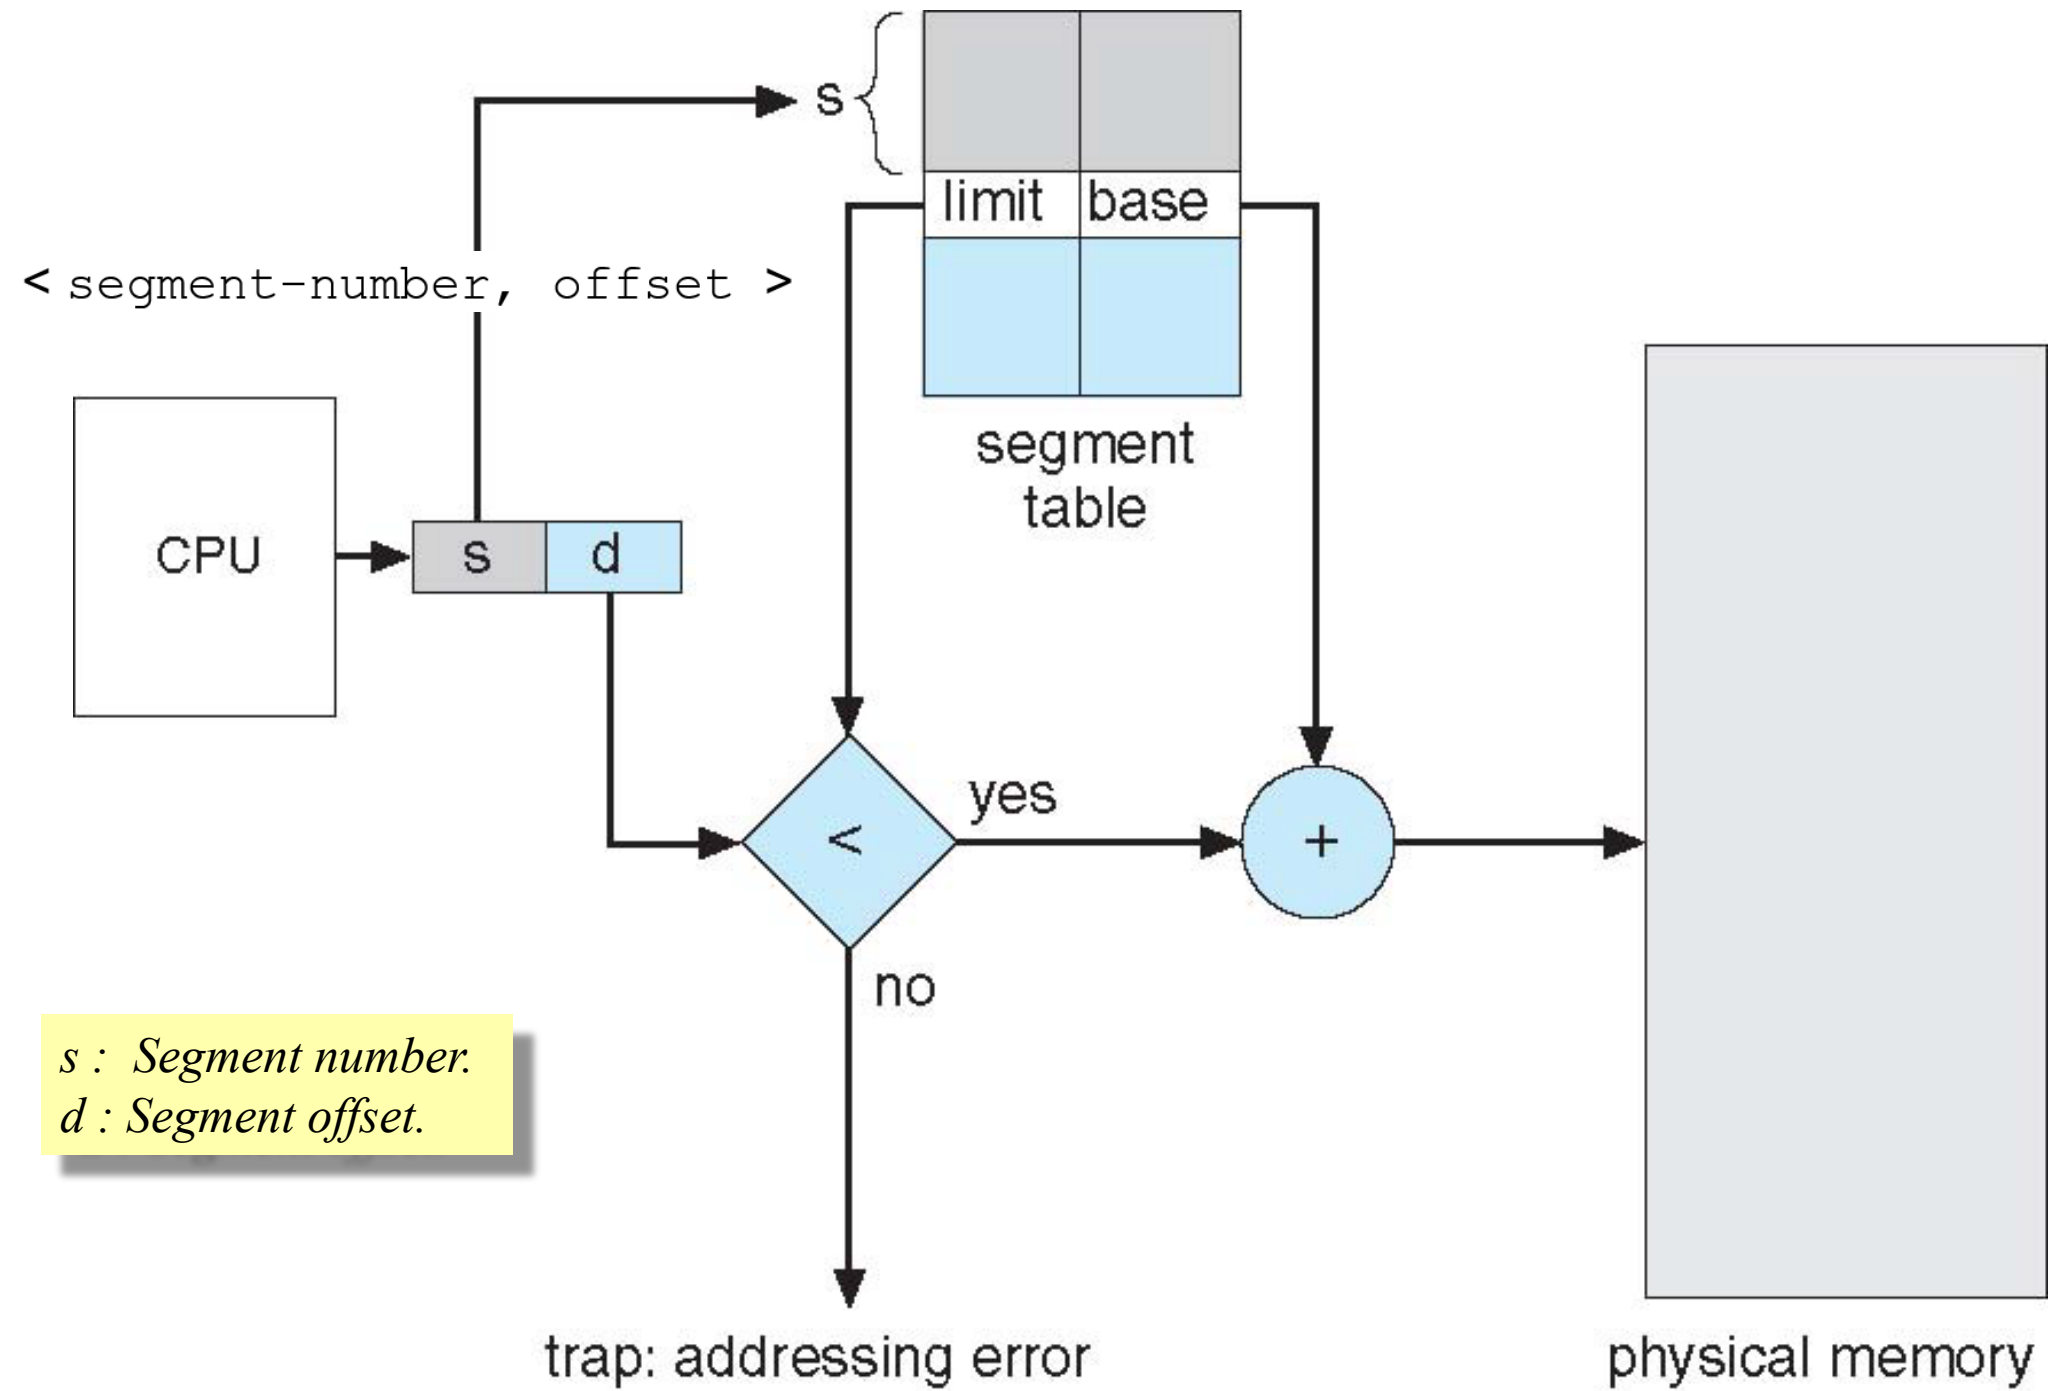
# Segmentation Hardware

- Addressing scheme require **segment number** and **displacement** / offset
- Logical address consists of a two tuple:

`< segment-number, offset >`

- Need to map a two-dimensional logical addresses to a one-dimensional physical address.
- Done via **segment table**:
  - **base** – contains the starting physical address where the segments reside in memory
  - **limit** – specifies the length of the segment.





**Segmentation hardware.**








# Example 9


What is the physical memory address for the logical address **2 : 53**?

< segment-number, offset >



|   | limit | base |
|---|-------|------|
| 0 | 1000  | 1400 |
| 1 | 400   | 6300 |
| 2 | 400   | 4300 |
| 3 | 1100  | 3200 |
| 4 | 1000  | 4700 |

segment table



Example of segmentation.





# Solution


What is the physical memory address for the logical address **2 : 53** ?

Segment 2 is 400 bytes long and begins at location 4300.

→ Physical address  
 $= 4300 + 53 = 4353$

|   | limit | base |
|---|-------|------|
| 0 | 1000  | 1400 |
| 1 | 400   | 6300 |
| 2 | 400   | 4300 |
| 3 | 1100  | 3200 |
| 4 | 1000  | 4700 |

segment table




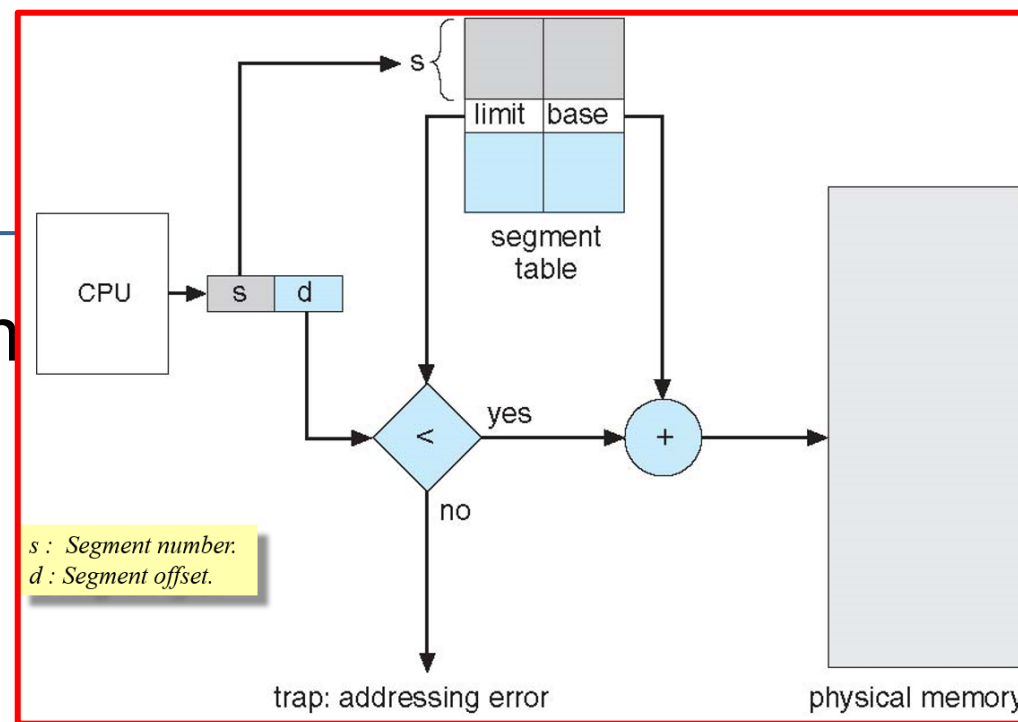
**4353**





What is the physical memory address?

- a) 3 : 852
- b) 0 : 1222



|   | limit | base |
|---|-------|------|
| 0 | 1000  | 1400 |
| 1 | 400   | 6300 |
| 2 | 400   | 4300 |
| 3 | 1100  | 3200 |
| 4 | 1000  | 4700 |

segment table

**Answer:**

**a) A reference to segment 3, byte 852, is mapped to Physical memory address:**

- = 3200 (the base of segment 3) + 852
- = 4052.

**b) A reference to byte 1222 of segment 0 would result in a trap to the OS, as this segment is only 1000 bytes long.**





## (b) Paging

- Physical address space of a process can be in **non-contiguous**.

- Process** is divided into **fixed-size** blocks, each of which may reside in a different part of physical memory.

1. Divide physical memory into **fixed-size** blocks, called **frames**:

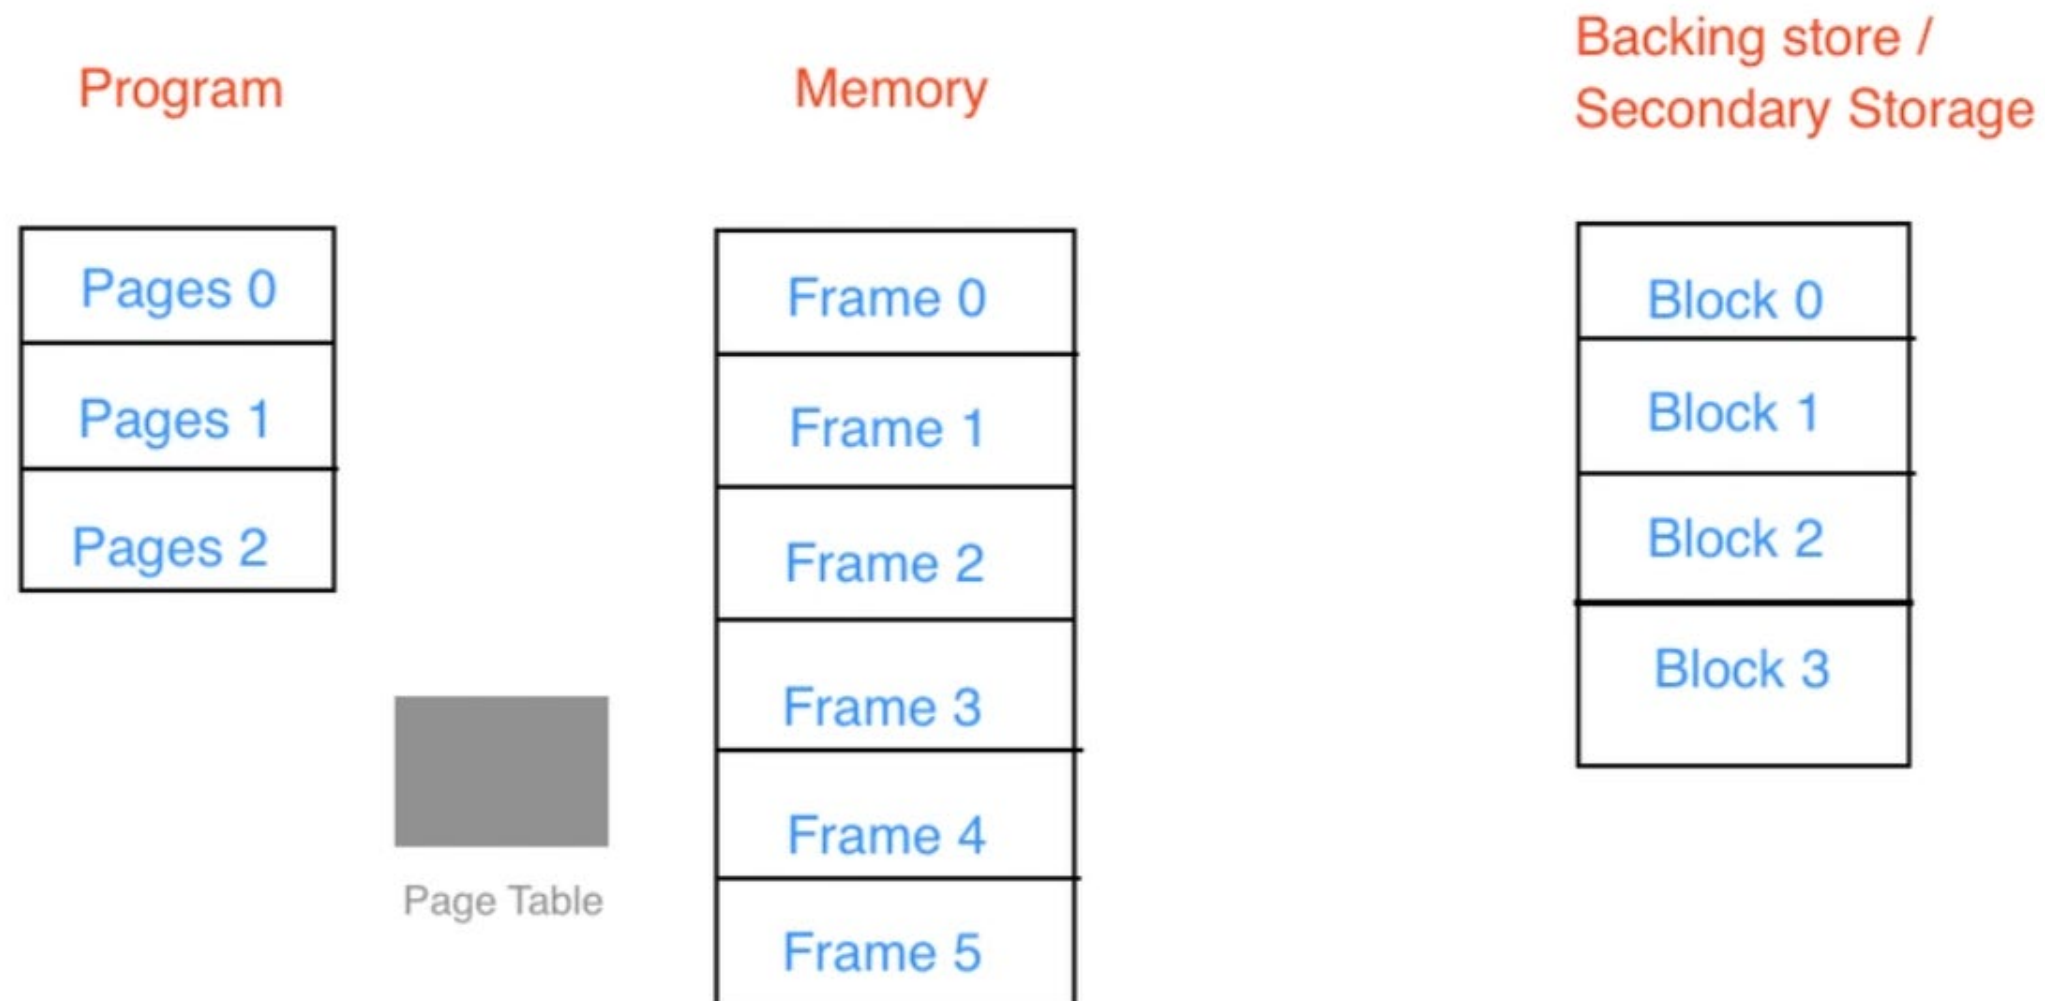
- Size of a frame is power of 2 between 512 Bytes and 8192 bytes ( $2^9 - 2^{13}$ ).

2. Divide logical memory into blocks of same size as frames, called **pages**.





3. Backing store, where the program is permanently residing, is also split into storage units called **blocks**, which are the same size as the **frames** and **pages**.






# Paging

Example 9a :

Consider a program with 350 lines long. Each page frame in a system can hold 100 lines. What happen in paged memory allocation?

Job 1





# Paging

## Solution 9a :

Job 1

|  |               |
|--|---------------|
|  | <i>Page 0</i> |
|  | <i>Page 1</i> |
|  | <i>Page 2</i> |
|  | <i>Page 3</i> |

Page size = Frame size

| Page Table |   |
|------------|---|
| 0          | 6 |
| 1          | 8 |
| 2          | 3 |
| 3          | 9 |

Page# → Frame#

*Internal  
Fragmentation* ----->

| Main Memory   | Page<br>Frame<br>No. |
|---------------|----------------------|
| OS            | 0                    |
|               | 1                    |
|               | 2                    |
| Job 1– Page 2 | 3                    |
|               | 4                    |
|               | 5                    |
| Job 1– Page 0 | 6                    |
|               | 7                    |
| Job 1– Page 1 | 8                    |
| Job 1– Page 3 | 9                    |
|               |                      |



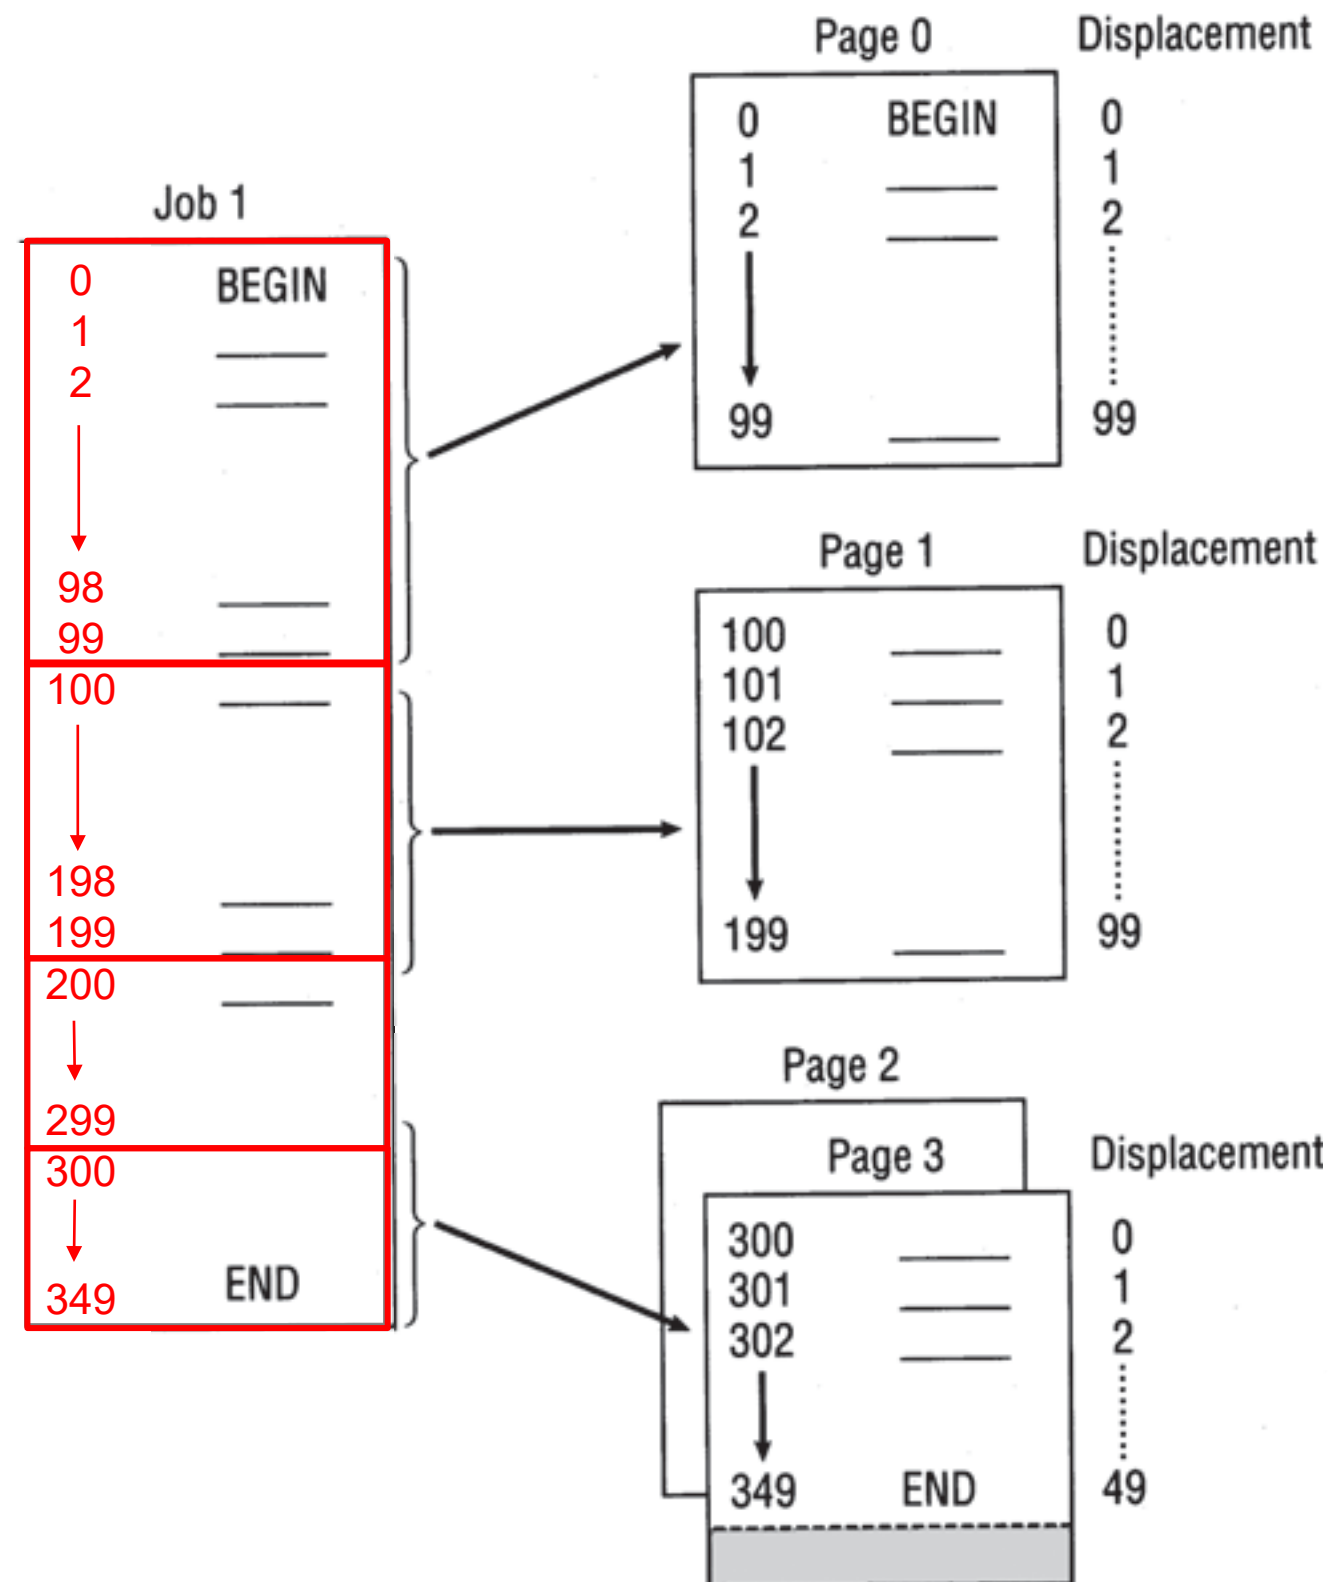




# Paging

## Example 9b :

- Job 1 is 350 lines long and is divided into 4 pages of 100 lines each.







- **Advantages:**
  - Allows job allocation in non-contiguous memory.
  - Efficient memory use.
- **Disadvantages:**
  - Increased overhead from address resolution.
  - **Internal fragmentation** in last page.
- Page size selection is crucial:
  - **Too small:** generates very **long page table**.
  - **Too large:** excessive **internal fragmentation**.






## Address translation scheme

- Address generated by CPU is divided into :
  - **Page number** ( $p$ ) – used as an index into a page table which contains base address of each page in physical memory.
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit.

Assume:

- ✓ logical address space =  $2^m$
- ✓ page size =  $2^n$





# Paging

Example 10 :


Considered size of logical address is 16 MB with it page size 1 KB. Find the bits for the page offset of the logical memory address.

Solution 10 :

- size of logical address = 16 MB =  $16 \times 2^{20}$  Bytes  
 $= 2^4 \times 2^{20} = 2^{24}$
- page size = 1 KB = 1024 Byte =  $2^{10}$

Page number

Page offset



$$\begin{aligned} &= m - n \\ &= 24 - 10 \\ &= 14 \text{ bits} \end{aligned}$$

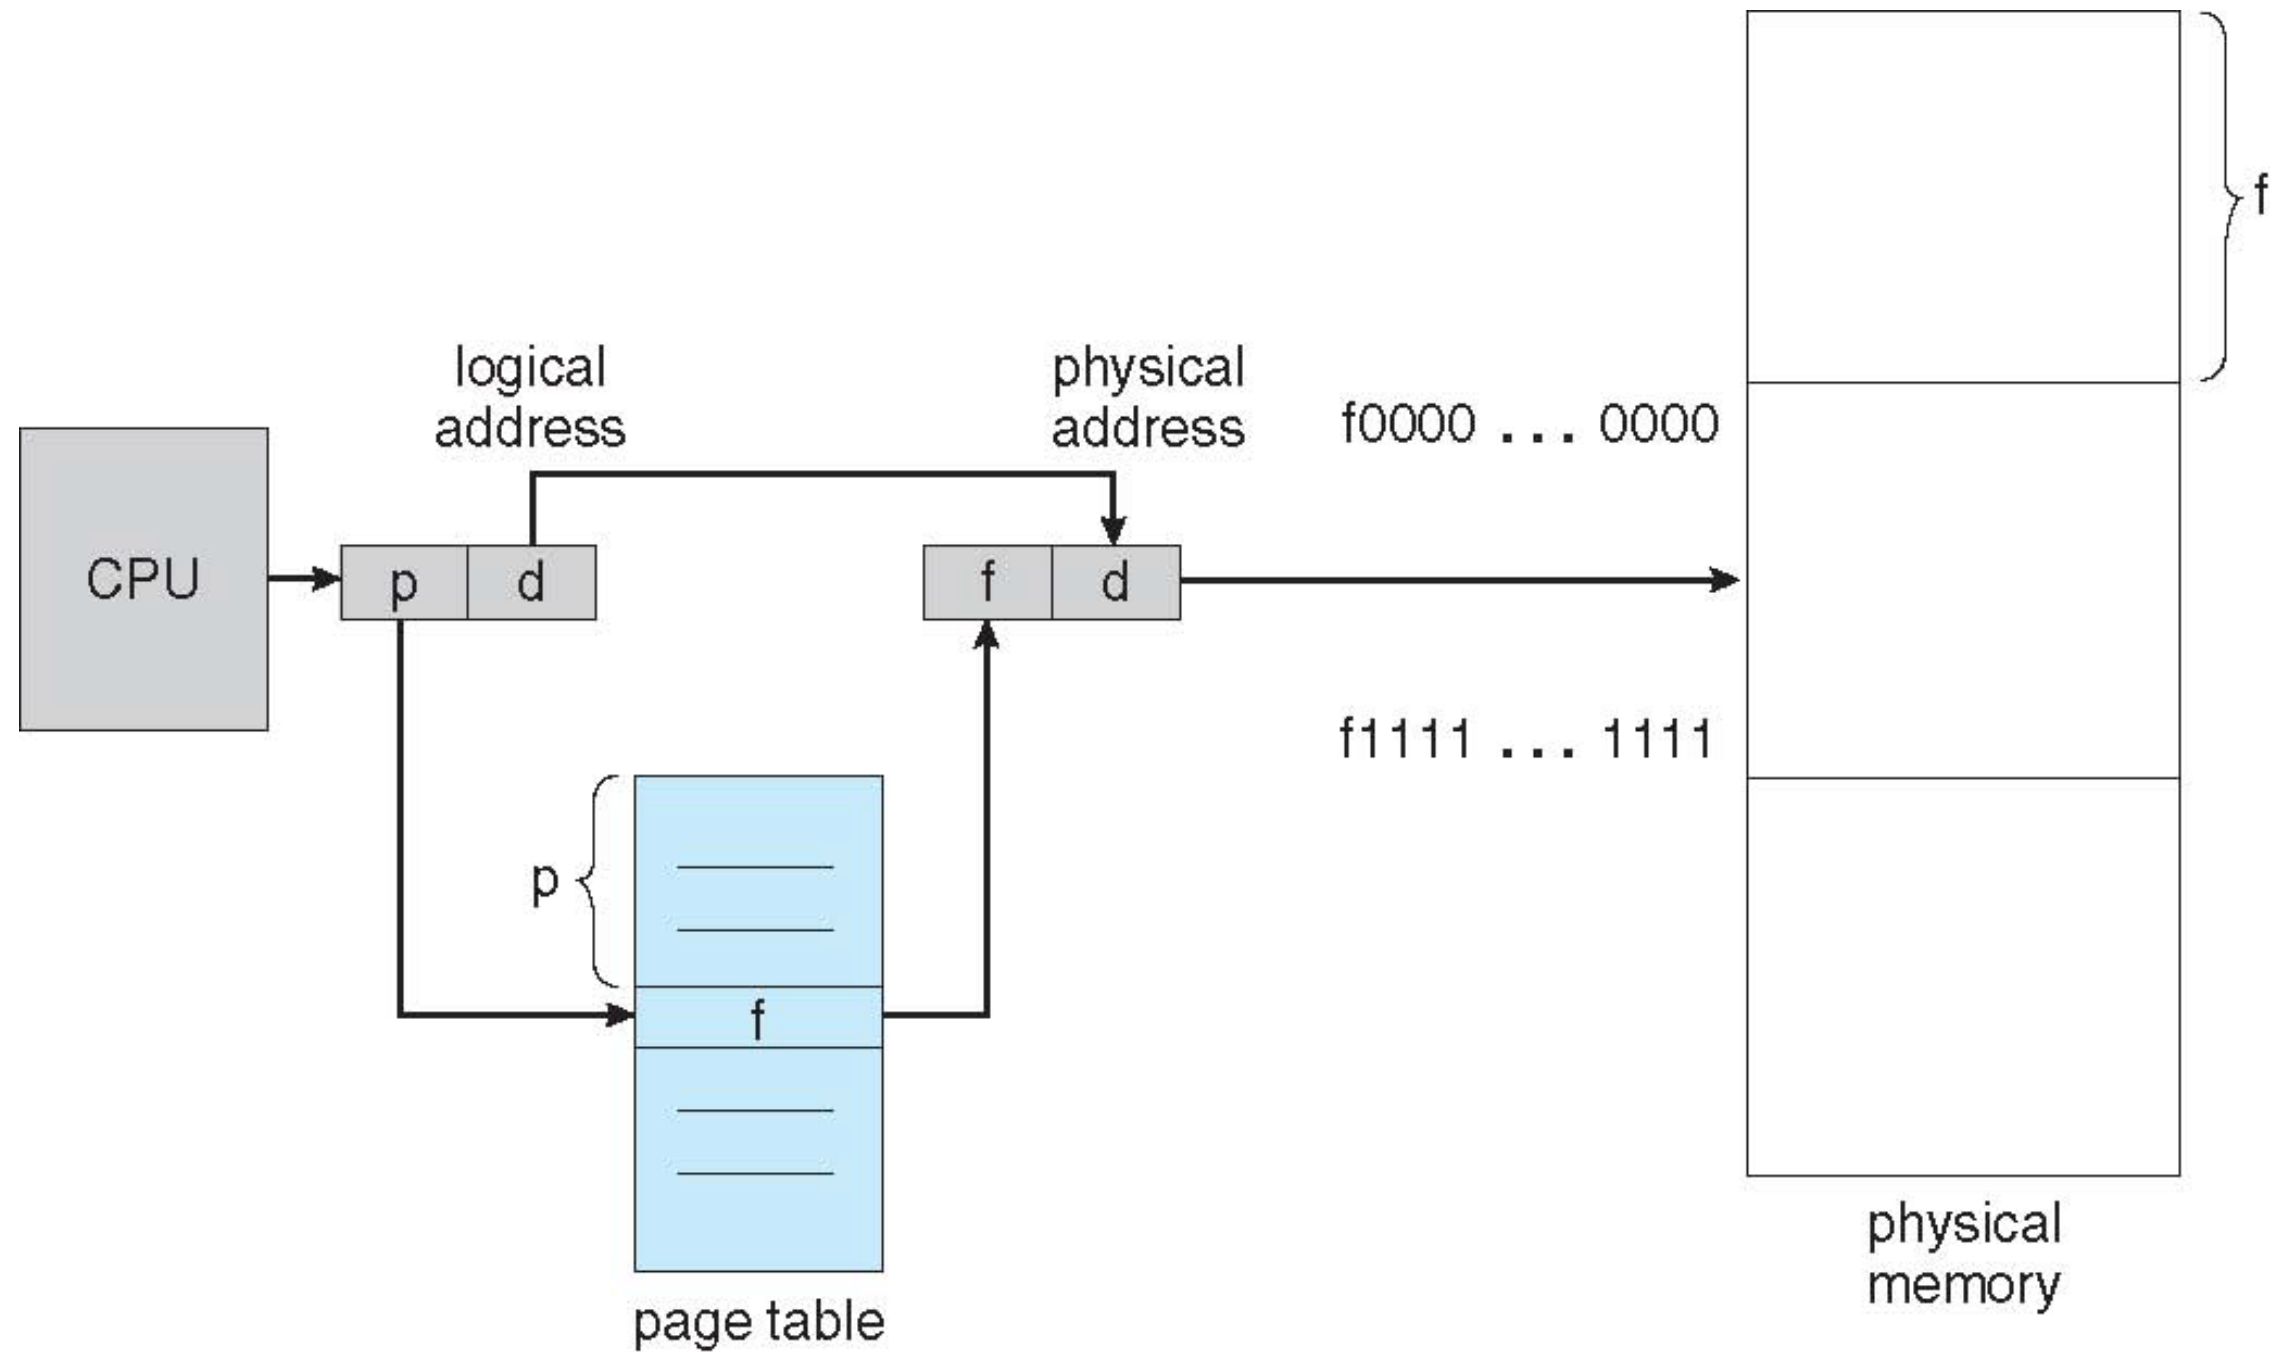
$$\begin{aligned} &= n \\ &= 10 \text{ bits} \end{aligned}$$





# Paging

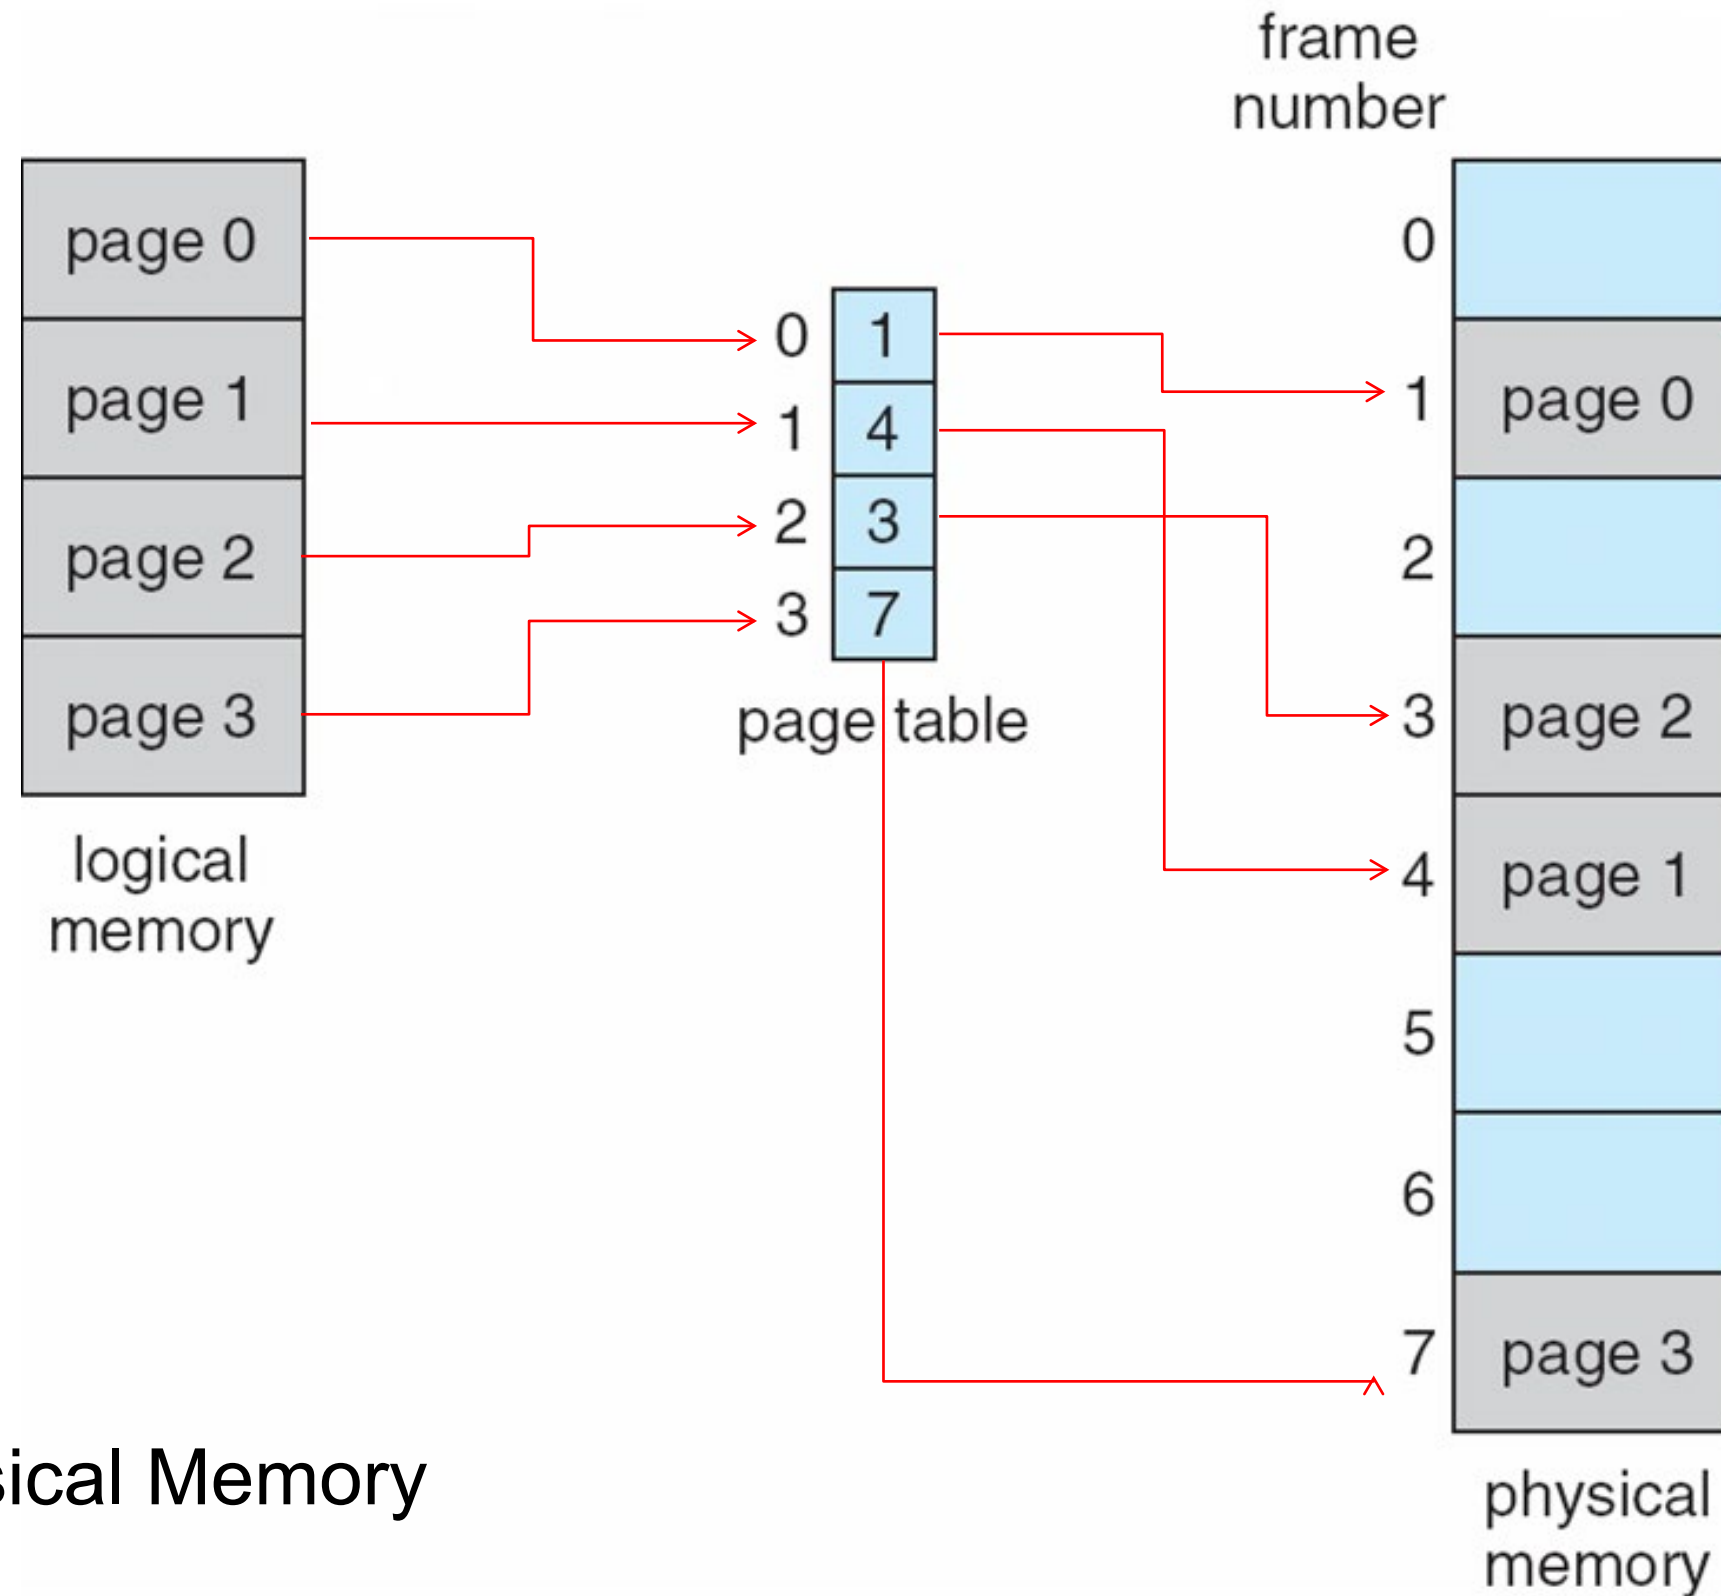
## Paging hardware





# Paging

## Example 11 :



Paging Model:  
Logical and Physical Memory





# Paging

## Example 12a :

Assume  $m = 4$   
and  $n = 2$  and  
32-byte memory  
and 4-byte  
pages.

|    |   |
|----|---|
| 0  | a |
| 1  | b |
| 2  | c |
| 3  | d |
| 4  | e |
| 5  | f |
| 6  | g |
| 7  | h |
| 8  | i |
| 9  | j |
| 10 | k |
| 11 | l |
| 12 | m |
| 13 | n |
| 14 | o |
| 15 | p |

logical memory

|   |   |
|---|---|
| 0 | 5 |
| 1 | 6 |
| 2 | 1 |
| 3 | 2 |

page table

|    |                  |
|----|------------------|
| 0  |                  |
| 4  | i<br>j<br>k<br>l |
| 8  | m<br>n<br>o<br>p |
| 12 |                  |
| 16 |                  |
| 20 | a<br>b<br>c<br>d |
| 24 | e<br>f<br>g<br>h |
| 28 |                  |

physical memory





# Paging

Page  
No.      Page  
Offset

|   |                  |
|---|------------------|
| 0 | 0<br>1<br>2<br>3 |
| 1 | 0<br>1<br>2<br>3 |
| 2 | 0<br>1<br>2<br>3 |
| 3 | 0<br>1<br>2<br>3 |

Page  
No:

0

1

2

3

|    |   |
|----|---|
| 0  | a |
| 1  | b |
| 2  | c |
| 3  | d |
| 4  | e |
| 5  | f |
| 6  | g |
| 7  | h |
| 8  | i |
| 9  | j |
| 10 | k |
| 11 | l |
| 12 | m |
| 13 | n |
| 14 | o |
| 15 | p |

logical memory

Page# → Frame#

|   |   |
|---|---|
| 0 | 5 |
| 1 | 6 |
| 2 | 1 |
| 3 | 2 |

page table

Frame  
No:

0

1

2

3

4

5

6

7

|    |                  |
|----|------------------|
| 0  |                  |
| 4  | i<br>j<br>k<br>l |
| 8  | m<br>n<br>o<br>p |
| 12 |                  |
| 16 |                  |
| 20 | a<br>b<br>c<br>d |
| 24 | e<br>f<br>g<br>h |
| 28 |                  |

physical memory

Frame  
No.      Frame  
Offset

|   |                  |
|---|------------------|
|   | 0<br>1<br>2<br>3 |
| 1 | 0<br>1<br>2<br>3 |
| 2 | 0<br>1<br>2<br>3 |
| 5 | 0<br>1<br>2<br>3 |
| 6 | 0<br>1<br>2<br>3 |





# Paging

## Example 12b :

Based on  
**Example 12a**,  
get the physical  
address for  $n$ .  
Show your work.

Page  
No:

0

1

2

3

|    |   |
|----|---|
| 0  | a |
| 1  | b |
| 2  | c |
| 3  | d |
| 4  | e |
| 5  | f |
| 6  | g |
| 7  | h |
| 8  | i |
| 9  | j |
| 10 | k |
| 11 | l |
| 12 | m |
| 13 | n |
| 14 | o |
| 15 | p |

logical memory

Page# → Frame#

|   |   |
|---|---|
| 0 | 5 |
| 1 | 6 |
| 2 | 1 |
| 3 | 2 |

page table

$n$

**Logical Address**

Page : Offset = 3 : 1

**Physical address**

$$\begin{aligned} &= ((\text{Frame No.}) \times (\text{Frame size})) \\ &\quad + \text{Offset} \\ &= (2 \times 4) + 1 \\ &= 8 + 1 = 9 \end{aligned}$$

Frame  
No:

0

1

2

3

4

5

6

7

|    |                  |
|----|------------------|
| 0  |                  |
| 4  | i<br>j<br>k<br>l |
| 8  | m<br>n<br>o<br>p |
| 12 |                  |
| 16 |                  |
| 20 | a<br>b<br>c<br>d |
| 24 | e<br>f<br>g<br>h |
| 28 |                  |

physical memory







# Paging

## Exercise 8.6 :

What is the physical address for the following logical addresses with given the number of the page and the offset?

a) 0 : 3

b) 1 : 0

c) 2 : 1

**Physical address**  
= ((Frame No.) x (Frame size))  
+ Offset

## Solution 8.6 :

a)

$$\begin{aligned} &= ( 5 \times 4 ) + 3 \\ &= 20 + 3 \\ &= \mathbf{23} \end{aligned}$$

b)

$$\begin{aligned} &= ( 6 \times 4 ) + 0 \\ &= 24 + 0 \\ &= \mathbf{24} \end{aligned}$$

c)

$$\begin{aligned} &= ( 1 \times 4 ) + 1 \\ &= 4 + 1 \\ &= \mathbf{5} \end{aligned}$$

|    |   |
|----|---|
| 0  | a |
| 1  | b |
| 2  | c |
| 3  | d |
| 4  | e |
| 5  | f |
| 6  | g |
| 7  | h |
| 8  | i |
| 9  | j |
| 10 | k |
| 11 | l |
| 12 | m |
| 13 | n |
| 14 | o |
| 15 | p |

logical memory

|   |   |
|---|---|
| 0 | 5 |
| 1 | 6 |
| 2 | 1 |
| 3 | 2 |

page table

|    |                  |
|----|------------------|
| 0  |                  |
| 4  | i<br>j<br>k<br>l |
| 8  | m<br>n<br>o<br>p |
| 12 |                  |
| 16 |                  |
| 20 | a<br>b<br>c<br>d |
| 24 | e<br>f<br>g<br>h |
| 28 |                  |

physical memory





**Exercise 8.7 :** Consider a paging system with 4 pages of logical address space, a page size of 4 bytes and a physical memory of 64 bytes as depicted in the following diagram. [8 marks]

**Logical Memory**

| Logical address | Content  |
|-----------------|----------|
| 0               | <b>A</b> |
| 1               | <b>B</b> |
| 2               | <b>C</b> |
| 3               | <b>D</b> |
| ·               | ·        |
| ·               | ·        |
| 8               | <b>G</b> |
| 9               | <b>H</b> |
| 10              | <b>I</b> |
| 11              | <b>J</b> |
| 12              | <b>K</b> |
| 13              | <b>L</b> |
| 14              | <b>M</b> |
| 15              | <b>N</b> |

- a) What is the frame size of the physical memory?
- b) How many frames of the physical memory?
- c) List the logical addresses for the following page.

| Page No. | Logical Address |
|----------|-----------------|
| 1        |                 |
| 3        |                 |





# Paging

d) Calculate the page offset for logical memory location that stores the data  $H$ .

d) Suppose a page table for the above paging system is as follows:

|   |    |
|---|----|
| 0 | 1  |
| 1 | 3  |
| 2 | 8  |
| 3 | 10 |

Calculate the physical address in binary for data  $H$ . [3 marks]





## Memory protection

- Memory protection implemented by associating **protection bits** with each frame to indicate if “*read-only*” or “*read-write*” access is allowed.
  - Can also add more bits to indicate “*execute-only*” and so on.
- **Valid-invalid** bit attached to each entry in the page table:
  - “*valid*” indicates that the associated page is in the process’ logical address space, and is thus is a legal page.
  - “*invalid*” indicates that the page is not in the process’ logical address space.
  - Or use **Page-Table Length Register (PTLR)**
- Any violations result in a trap to the kernel.





# Paging

## Example 13 :

(Page: 370)




Figure 8.15 Valid (v) or Invalid (i) Bit In A Page Table.





# Summary





# Summary

- Memory-management algorithms for multiprogrammed operating systems range from the simple single-user system approach to segmentation and paging.
- The most important determinant of the method used is the **hardware** provided.
  - ✓ Every memory address generated by the CPU must be **checked** for legality and possibly mapped to a physical address.
  - ✓ Cannot be implemented in software.



# End of Chapter 8

---

